

Lecture 11

Sorting 2

EECS-214

Why is sorting hard?


- Have to **compare a lot** of elements
- Naïve algorithms involve **exhaustive comparisons**
 - **Every element** to every other element
 - $O(n^2)$

Question

- What can we do by comparing a **single element** to every other element?
 - $O(n)$


Partitioning

- Comparing one element to all the others lets us **divide the array** into the elements
 - **less than**,
 - **greater than**,
 - and **equal to** the original element
- This is called **partitioning** the array
 - And the element being compared to is called the **pivot** element



Partitioning


- Partitioning makes the array “**more sorted**”
- But we’re not done yet
 - The left- and right- sides **aren’t sorted internally**
- What can we do to finish sorting?



Partitioning

Recurse!

- Use partitioning to make the left and right sides more sorted too
- Keep going until everything is fully sorted



Quicksort (Hoare, 1960)

- One of the most famous algorithms in computer science
- **Approximately sorts** whole array
- **Recurse**s to complete sorting of each half

Quicksort(array)

select pivot element

partition about pivot

Quicksort(left side)

Quicksort(right side)

Quicksort (Hoare, 1960)

- This code is pretty **hand-wavy**
 - For example what does it mean to call quicksort of the left side of the array?
 - Also, there's nothing to stop the recursion
- Making it more **rigorous** is **straightforward**
 - But there are more confusing details
 - So it's good to understand the basic idea first

Quicksort(array)

select pivot element

partition about pivot

Quicksort(left side)

Quicksort(right side)

Implementing partitioning

Basic idea

- **Move pivot** to the **right** (the end)
- **Sweep from left** to right
 - If you see an element that's **less than** the pivot
 - **Swap** it with something on the left that's bigger than the pivot
- **Move pivot between** the less-than and greater-than elements

```
Partition(a, pIndex)
    pValue = a[pIndex]
    // Move pivot to end
    swap a[pIndex]
        with last element of a
    // Move small values left
    nextLeft = 0
    for i = 0 to a.Length-1
        if a[i] <= pValue
            swap a[i]
                with a[nextLeft]
            nextLeft++
    // Move the pivot into place
    swap a[nextLeft]
        with last element of a
```


Implementing partitioning

This algorithm is a little inscrutable

- This is an algorithm that induces a kind of behavior in textbook writers that I hate: the **pseudoexplanation**

```
Partition(a, pIndex)
    pValue = a[pIndex]
    // Move pivot to end
    swap a[pIndex]
        with last element of a
    // Move small values left
    nextLeft = 0
    for i = 0 to a.Length-1
        if a[i] <= pValue
            swap a[i]
                with a[nextLeft]
            nextLeft++
    // Move the pivot into place
    swap a[nextLeft]
        with last element of a
```

Implementing partitioning

Wikipedia's pseudoexplanation:

- “This is the in-place partition algorithm. It partitions ... the array ... by moving all elements less than or equal to [the pivot value] to the beginning of the subarray, leaving all the greater elements following them”
- This is simply a **restatement** of what the procedure **should do**
- It's not an explanation of **why it works**

```
Partition(a, pIndex)
    pValue = a[pIndex]
    // Move pivot to end
    swap a[pIndex]
        with last element of a
    // Move small values left
    nextLeft = 0
    for i = 0 to a.Length-1
        if a[i] <= pValue
            swap a[i]
                with a[nextLeft]
            nextLeft++
    // Move the pivot into place
    swap a[nextLeft]
        with last element of a
```

Implementing partitioning


Okay, why does it actually work?

```
Partition(a, pIndex)
    pValue = a[pIndex]
    // Move pivot to end
    swap a[pIndex]
        with last element of a
    // Move small values left
    nextLeft = 0
    for i = 0 to a.Length-1
        if a[i] <= pValue
            swap a[i]
                with a[nextLeft]
            nextLeft++
    // Move the pivot into place
    swap a[nextLeft]
        with last element of a
```

Implementing partitioning

Claim: the loop maintains the **invariants** that:

- Elements **before nextLeft** are \leq **the pivot**
- Elements **between nextLeft and i**, but not including i, are $>$ **the pivot**




Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Proof:


- For the first iteration,
 $i = \text{nextLeft} = 0$
- So there are no $j < \text{nextLeft}$
or $j < i$
- So it's trivially true

Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Proof:


- Now assume it's true at the end of iteration $i-1$
- Let's show it's true at the end of iteration i

Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Proof:


- So we know
 - $a[j] \leq \text{pivot}$, for $j < \text{nextLeft}$
 - $a[j] > \text{pivot}$, for $\text{nextLeft} \leq j < i-1$
- Case 1: $a[i] > \text{pivot}$
 - Don't do anything
 - So $a[j] \leq \text{pivot}$, for $j < \text{nextLeft}$
 - $a[j] \leq \text{pivot}$, for $\text{nextLeft} \leq j < i-1$
 - but $a[i] \text{ also } > \text{pivot}$, so
 - $a[j] > \text{pivot}$, for $\text{nextLeft} \leq j < i$

Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Proof:


- So we know
 - $a[j] \leq \text{pivot}$, for $j < \text{nextLeft}$
 - $a[j] > \text{pivot}$, for $\text{nextLeft} \leq j < i-1$
- Case 2: $a[i] \leq \text{pivot}$
 - We swap $a[i]$ and $a[\text{nextLeft}]$
 - $a[\text{nextLeft}]$ is now $\leq \text{pivot}$
 - And we increment nextLeft
 - So once again, $a[j] \leq \text{pivot}$, for $j < \text{nextLeft}$

Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Proof:


- So we know
 - $a[j] \leq \text{pivot}$, for $j < \text{nextLeft}$
 - $a[j] > \text{pivot}$, for $\text{nextLeft} \leq j < i-1$
- Case 2: $a[i] \leq \text{pivot}$
 - We swap $a[i]$ and $a[\text{nextLeft}]$
 - This also means $a[i]$ is now $> \text{pivot}$
 - $a[j]$ still $> \text{pivot}$, for $\text{nextLeft} \leq j < i-1$
 - but $a[i]$ also $> \text{pivot}$, so
 - $a[j] > \text{pivot}$, for $\text{nextLeft} \leq j < i$

Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Proof:


- So by induction, the theorem holds for each iteration

Implementing partitioning

Theorem: on each iteration of the for loop:

For any j :

- $a[j] \leq \text{pivot}$, if $j < \text{nextLeft}$
- $a[j] > \text{pivot}$, if $\text{nextLeft} \leq j < i$



Corollary: the stupid algorithm works

- We run until i gets up to the end (where the pivot element is)
- Everything before nextLeft is $\leq \text{pivot}$
- Everything from nextLeft to pivot is $> \text{pivot}$
- Then we swap nextLeft (which is $> \text{pivot}$) with the pivot
- So we end up with:
 - Before nextLeft is $\leq \text{pivot}$
 - nextLeft is the pivot
 - After nextLeft is $> \text{the pivot}$
- Yay!


In-place partitioning

9	3	7	4	8
---	---	---	---	---

Pivot around 7

- pIndex = 2
- pValue = 7

In-place partitioning




Swap pivot to end

Pivot around 7

- pIndex = 2
- pValue = 7

In-place partitioning




↑
↑ i
| nextLeft

start i, nextLeft
at the beginning


Pivot around 7

- pValue = 7

In-place partitioning




$a[i] \leq \text{pValue}?$



Pivot around 7

- $\text{pValue} = 7$

In-place partitioning




no – move on

↑
↑ i
| nextLeft

Pivot around 7

- pValue = 7

In-place partitioning




$a[i] \leq \text{pValue}?$

Pivot around 7

- $\text{pValue} = 7$

In-place partitioning



↑
nextLeft


↑ i

yes – swap with
a[nextLeft]


Pivot around 7

- pValue = 7

In-place partitioning




increment nextLeft



Pivot around 7

- pValue = 7


In-place partitioning



Pivot around 7

- pValue = 7

In-place partitioning




$a[i] < \text{pValue}$?

Pivot around 7

- $\text{pValue} = 7$

In-place partitioning




no – move on

Pivot around 7

- pValue = 7

In-place partitioning




$a[i] \leq \text{pValue}$?

Pivot around 7

- $\text{pValue} = 7$

In-place partitioning




yes - swap

Pivot around 7

- pValue = 7

In-place partitioning




increment nextLeft

Pivot around 7

- pValue = 7

In-place partitioning

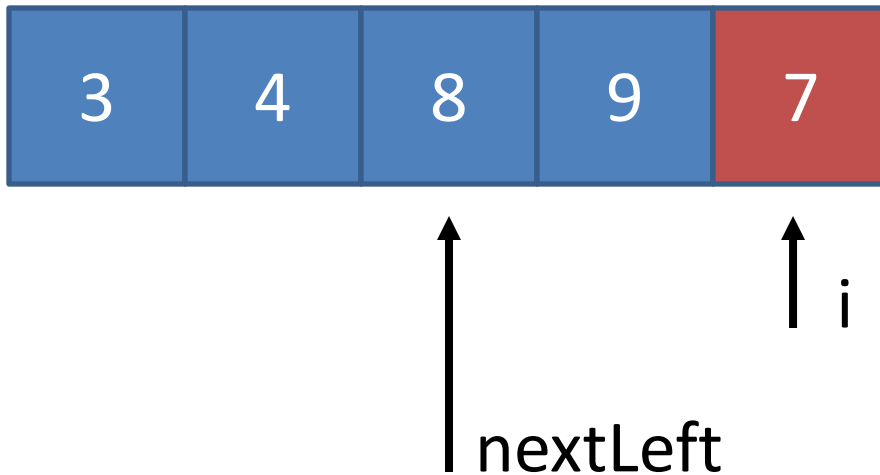


move on

Pivot around 7

- pValue = 7

In-place partitioning

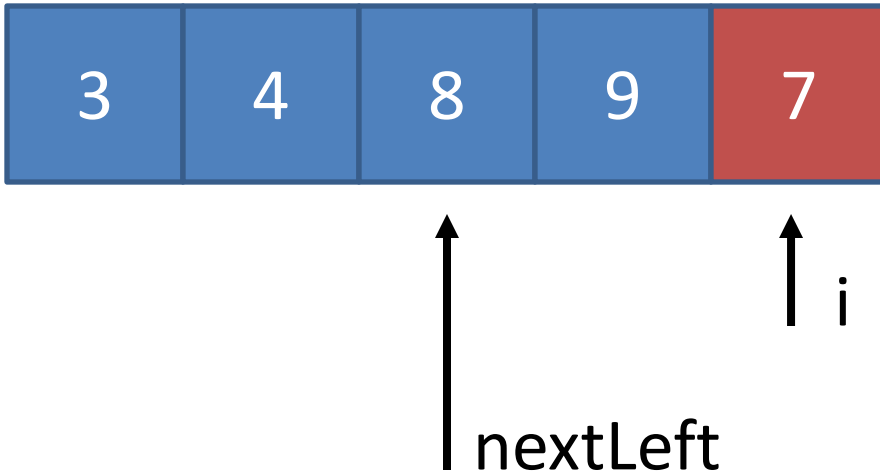


oops –
we're at the pivot

Pivot around 7

- pValue = 7

In-place partitioning




swap it into nextLeft

Pivot around 7

- pValue = 7

In-place partitioning



done

Pivot around 7

- pValue = 7

Partitioning subranges of the array

- When quicksort recurses, it will need to partition **just the left and right halves** of the array
- So we need to change partition to take the **region of the array** to be partitioned as an **argument**
 - Specified by its starting and ending indices
- Quicksort also needs Partition to tell it where the split is, so it will return nextLeft

```
Partition(a, pIndex, start, end)
    pValue = a[pIndex]
    // Move pivot to end
    swap a[pIndex]
        with a[end]
    // Move small values left
    nextLeft = start
    for i = start to end-1
        if a[i] <= pValue
            swap a[i]
                with a[nextLeft]
            nextLeft++
    // Move the pivot into place
    swap a[nextLeft]
        with a[end]
    return nextLeft
```

The real quicksort algorithm

```
Quicksort(a)
```

```
    QuicksortRecur(a, 0, a.Length-1)
```

```
QuicksortRecur(a, start, end)
```

```
    if end > start
```

```
        pIndex = ... select pivot index ...
```

```
        newIndex = Partition(a, pIndex, start, end)
```

```
        QuicksortRecur(a, start, newIndex-1)
```

```
        QuicksortRecur(a, newIndex+1, end)
```

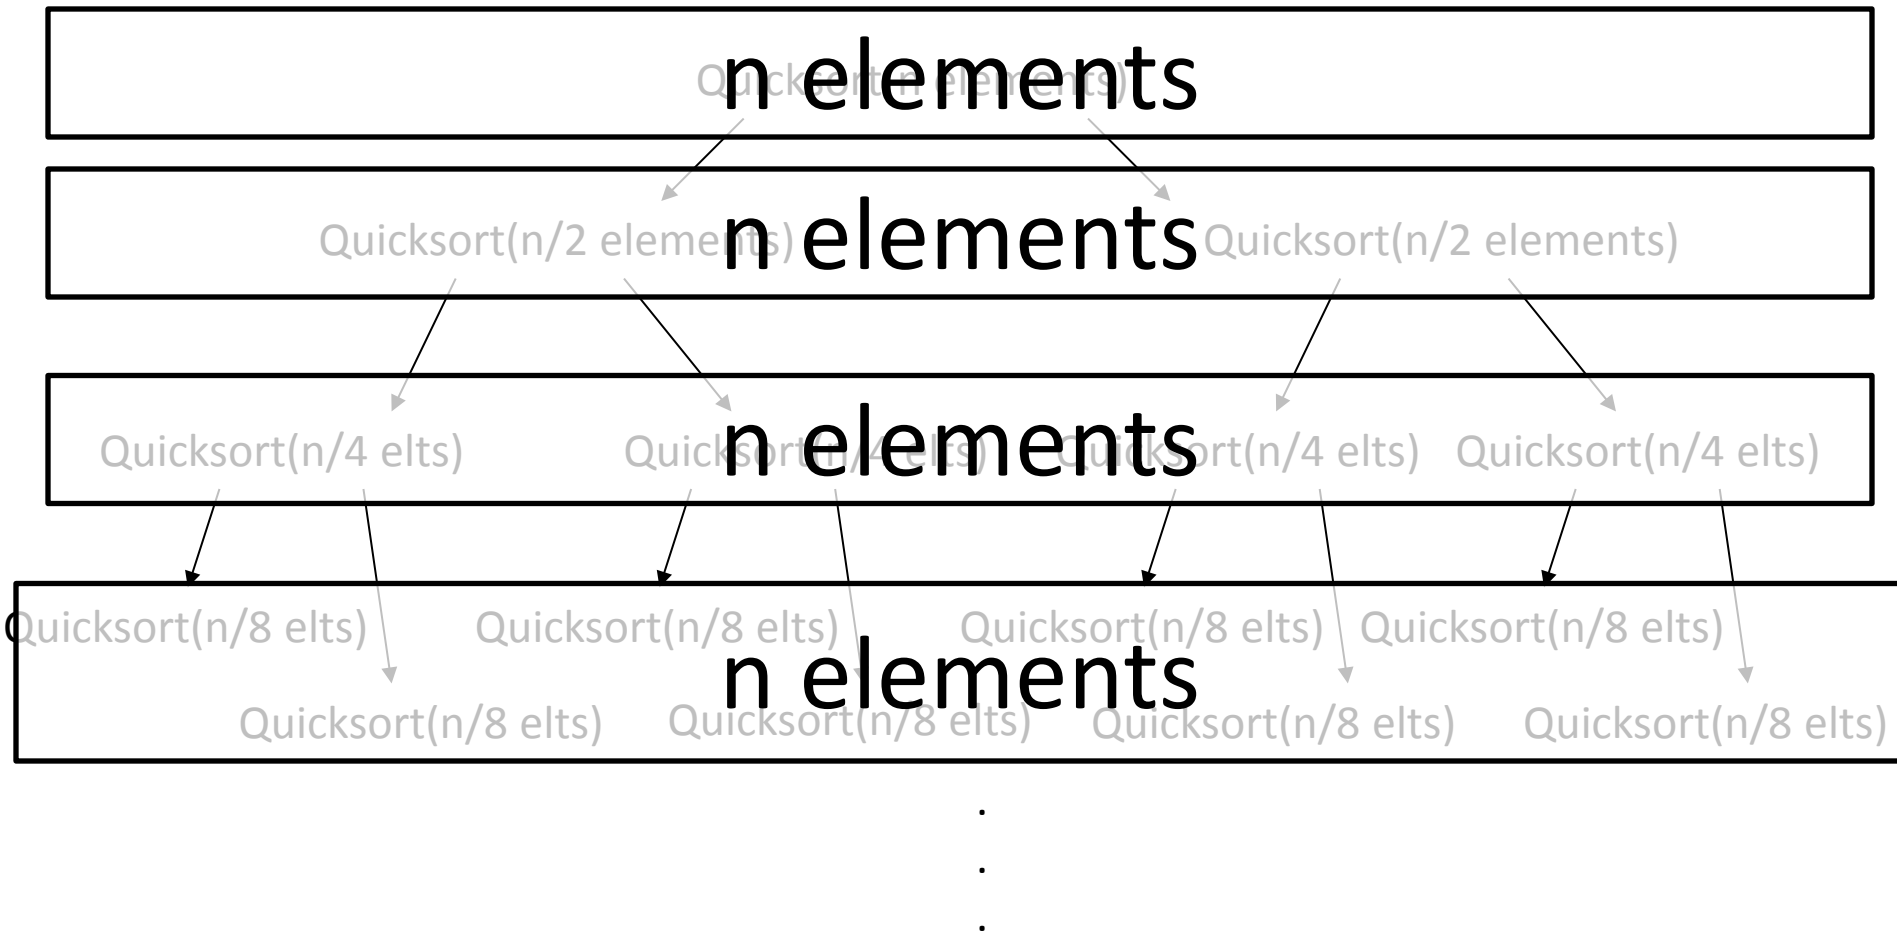
We'll talk about
this part shortly ...



Analysis (best case)

- The best case is where we **always pick** a pivot that splits the array **exactly in half**
- Then we recurse **$\lceil \log_2 n \rceil$** times
 - $\lceil x \rceil$ means “ x rounded up to the nearest integer
 - Pronounced “**ceiling** x ”
 - $\lfloor x \rfloor$ means round down (read: “**floor** x ”)

Call tree



Note: this is somewhat hand-wavy, since each level of recursion removes one element (the pivot) before dividing by two. But this still works as an upper bound

Analysis (best case)

- $O(\log n)$ levels of recursion
- n items total partitioned in each level
- **$O(n \log n)$** total time

Analysis (worst case)

- The worst case is where we always **pick largest or smallest** element as the pivot
 - All remaining elements on one side
- Then we recurse **$n - 1$** times!
 - i.e. n total calls

Analysis (worst case)

- $O(n)$ levels of recursion
- $O(n)$ work per level of recursion
- $O(n^2)$ total time

Analysis (average case)

Hand-wavy argument:

- You **have to pick really badly** to get worst-case behavior
- Even if you only split it so that each time:
 - 99% of the elements are on one side
 - 1% are on the the other side
 - You still get $O(\log n)$ levels of recursion
 - It's just closer to $7 \log n$ than $\log n$, but that's still $O(\log n)$
- So even if most of the time you only get the split to **within 1%, you still get $O(n \log n)$** total execution time

Picking the pivot element

- So the choice of the pivot **determines performance**
- How do we choose the pivot element?

Technique 1: fixed choice

- E.g. always pick the **first or last** element in the array
- Usually works fine, but can give you worst-case behavior if the array is already sorted/reverse-sorted
 - Example: first or last element gives you worst-case behavior on **both sorted**, and **reverse-sorted** arrays

Technique 2: median of 3

- Grab **three elements** (e.g. first, last, and middle of the array)
- Find the **median** (the one that's between the other two)
- Use that as the pivot
- Works well, but a **sneaky adversary** could cook up an array to force your sort to run slowly
 - Mostly of theoretical interest

Technique 3:

Median of the whole array

- Compute the median (middle element) of the whole array
 - Unfortunately, simple methods for computing the media are $O(n^2)$ in the worst case
 - There are $O(n)$ **average case** methods, but they have **large constant factors**
- This is called **median sort**

Technique 4: choose randomly

- This is what people generally use in practice
- $O(n \log n)$ behavior on average
- Doesn't depend on the input
 - So there's **no worst-case input**
 - Just **worst-case luck**

Reading

- Read CLR chapter on Quicksort (chapter 8/9, depending on edition)