

Lecture 13

Heaps and priority queues

EECS-214

Queues

- Simplified **sequence** data structure
 - **Insertions** only at the end (“**tail**”)
 - **Deletions** only from the beginning (“**head**”)
- **First-in, first out**
 - Objects are dequeued in the order they were enqueued
- Simple API
 - **Enqueue**: add item to the end of the queue
 - **Dequeue**: remove item from the front



Priority queues

- Like normal queues
 - Objects **wait in line** to be processed
- However, items have an associated **numeric priority**
 - Priority specified when added to queue
 - Objects removed from queue in order of priority
- Slightly different API
 - **Insert**(object, priority)
 - Adds object with specified priority
 - **ExtractMax**()
 - Returns highest priority object



Priority queues

- Like normal queues
 - Objects wait in line to be processed
- However, items have an associated numeric priority
 - Priority specified when added to queue
 - Objects removed from queue in order of priority
- Slightly different API
 - Insert(object, priority)
 - Adds object with specified priority
 - **ExtractMin()**
 - Returns **lowest** priority object



Implementing priority queues

- We can use a **balanced tree** (e.g. a red/black tree) as a priority queue

- Insert using the normal RBT insert

- $O(\log n)$ time

- Extract max is


- Get the maximum element

- $O(\log n)$

- Delete it


- Also $O(\log n)$

- **$O(\log n)$** time




Implementing priority queues

- Unfortunately, red/black trees are **pretty complicated**
 - They go to a lot of work to keep all the items perfectly sorted
- Is there something **simpler** we could do?




Heaps



- Heaps are a simple **tree structure** for implementing **priority queues**
- Rather than requiring their in-order traversal to be sorted
 - We just require that **parent nodes be larger than** their child nodes
 - Or **smaller**, if it's a **min heap**
- There are lots of exotic types of heaps
 - We'll focus on binary heaps
 - Which are **complete binary trees** with the heap property
 - We'll get to the completeness thing in a minute...

Heaps



Proposition: the **largest element** of a heap is always its **root**


Proof:

- Suppose some **other element** is the largest element
- Since it isn't the root, it **must have a parent**
- Since it's the largest element, it **must be larger than its parent**
- But that **contradicts** the definition of a heap
- So the largest element must be the root

Sketch of heap algorithms

ExtractMax (version 1)


- We know the **root** is the **maximal** element
- So we want to **delete it** and **return it**
- But we need to **replace it** with its **largest child**
 - So we find the largest child
 - And **recursively delete it** from its subtree



Sketch of heap algorithms

ExtractMax (version 1)

- We know the root is the maximal element
- So we want to delete it and return it
- But we need to replace it with its largest child
 - So we find the largest child
 - And recursively delete it from its subtree




However, this is going to turn out to not to be the most convenient algorithm

Sketch of heap algorithms

ExtractMax (version 2)


- Replace the root with a **leaf node**
 - Getting to a leaf will **turn out to be easy**
- The leaf node is **probably too small**
- So **move it downward** in the tree to make it be a proper heap again



Sketch of heap algorithms

ExtractMax(heap)

result = root of heap
replace root with leaf
Heapify(heap)
return result



Sketch of heap algorithms

Heapify(root)

if root > children

done

else if left child > root and
left child > right child


swap root and left child

Heapify(left child)

else


swap root and right child

Heapify(right child)




Example

- Starting configuration




Example

- **Replace root** with **leaf**
 - **Violates** heap property




Example

- Replace root with leaf
- **Swap** with **largest** child
 - Still violates heap property




Example

- Replace root with leaf
- Swap with largest child
- **Swap** with **largest** child again
 - Still violates heap property



Example


- Replace root with leaf
- Swap with largest child
- Swap with largest child again
- **Swap** with **largest** child again
 - Now a **proper heap**



Sketch of heap algorithms


Insert(item)

- **Add item as a leaf**
 - Again, trust us, this will turn out to be easy
- While item > its parent
 - **Swap with parent**
 - Compare it to the next level up




Analysis

- Both algorithms
 - Move nodes up/down tree
 - Perform a **constant** amount of work **at each level**
- So their execution time is **$O(h)$**
 - Where h is the tree's height
- Again, this is **good, if the tree is balanced**, bad otherwise




Complete binary trees

- A **complete** binary tree is a binary tree in which
 - **Every level** of the tree is **full**, **except** possibly the **last**
 - Can't add anymore nodes
 - Every node is **shifted** as far to the **left** as possible
- Complete trees are **optimally balanced**




Binary heaps

- A **binary heap** is a
 - Complete binary tree
 - That satisfies the heap property
- Great!
- How do we ensure that the heap is a complete binary tree?




Embedding in an array

- It turns out that any complete binary tree can be **embedded** an array in a particularly cleaver way
- We can **compute**
 - The position of its **parent** in the array,
 - and the positions of its **children**,
 - directly from its own **position**





Embedding in an array

- Store the **root** in the first element (**element 0**)
- For any node
 - Let i be its position in the array (for the root, $i = 0$)
 - Store its **left child** at position $2i + 1$
 - Store its **right child** at position $2i + 2$
 - Its parent can be found at position $\lfloor (i - 1)/2 \rfloor$
- Trust me that this works :-)



Why is this a good representation?

- Very **fast**
 - Can just **allocate a big array** and then never have to call new again
- **Last element** is always a **leaf**
 - Remember our algorithms needed to **add/remove leaves**?





Representing a heap using an array

- Assume we have an extra field for the array to keep track of the size of the heap
- Define the following utility procedures:

Parent(int i)
 return $(i-1)/2$

Left(int i)
 return $2*i+1$

Right(int i)
 return $2*i+2$



Heap insertion using the array representation

HeapInsert(A, key)

$A.size = A.size + 1$

$i = A.size$



while $i > 0$ and

$A[Parent(i)] < key$

$A[i] = A[Parent(i)]$

$i = Parent(i)$

$A[i] = key$



Inserting 10

HeapInsert(A, key)

A.size = A.size + 1

i = A.size



while i > 0 and

$A[\text{Parent}(i)] < \text{key}$

$A[i] = A[\text{Parent}(i)]$

$i = \text{Parent}(i)$

$A[i] = \text{key}$



Check parent

HeapInsert(A, key)

A.size = A.size + 1

i = A.size



while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



$$\text{Parent}(i) = (5-1)/2 = 2$$

$8 < 10$, so $A[\text{Parent}(i)] < \text{key}$

HeapInsert(A, key)

$A.\text{size} = A.\text{size} + 1$

$i = A.\text{size}$



while $i > 0$ and

$A[\text{Parent}(i)] < \text{key}$

$A[i] = A[\text{Parent}(i)]$

$i = \text{Parent}(i)$

$A[i] = \text{key}$



$$\text{Parent}(i) = (5-1)/2 = 2$$

Copy parent down

HeapInsert(A, key)

A.size = A.size + 1

i = A.size



while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



$$\text{Parent}(i) = (5-1)/2 = 2$$

And move up tree

HeapInsert(A, key)

A.size = A.size + 1

i = A.size


while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



Check parent

HeapInsert(A, key)

A.size = A.size + 1

i = A.size



while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



$\text{Parent}(i) = (2-1)/2 = 0$
(remember int arithmetic rounds down)

$9 < 10$, so $A[\text{Parent}(i)] < \text{key}$

HeapInsert(A, key)

$A.\text{size} = A.\text{size} + 1$

$i = A.\text{size}$



while $i > 0$ and

$A[\text{Parent}(i)] < \text{key}$

$A[i] = A[\text{Parent}(i)]$

$i = \text{Parent}(i)$

$A[i] = \text{key}$



$\text{Parent}(i) = (2-1)/2 = 0$
(remember int arithmetic rounds down)

Copy parent down

HeapInsert(A, key)

A.size = A.size + 1

i = A.size


while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



i=2

Parent(i) = (2-1)/2 = 0
(remember int arithmetic rounds down)

Move up

HeapInsert(A, key)

A.size = A.size + 1

i = A.size



while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



Can't move farther

HeapInsert(A, key)

$A.size = A.size + 1$

$i = A.size$



while **$i > 0$** and

$A[Parent(i)] < key$

$A[i] = A[Parent(i)]$

$i = Parent(i)$

$A[i] = key$



Store the new key

HeapInsert(A, key)

$A.size = A.size + 1$

$i = A.size$



while $i > 0$ and

$A[Parent(i)] < key$

$A[i] = A[Parent(i)]$

$i = Parent(i)$

$A[i] = key$



Done!

HeapInsert(A, key)

A.size = A.size + 1

i = A.size


while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



Notice that this is once again a valid heap

Extracting an element

HeapExtractMax(A)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Extracting an element

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i


if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



Here we go!

HeapExtractMax(A)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Remember the max (the root)

HeapExtractMax(A)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Move the last leaf to the root

HeapExtractMax(A)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Move the last leaf to the root

HeapExtractMax(A)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Re-heapify

HeapExtractMax(A)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Re-heapify

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i


if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap $A[i]$ and $A[largest]$

Heapify(A, largest)



Find the left- and right-children

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



$A[l] \text{ not } > A[i]$

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.\text{size}$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.\text{size}$ and $A[r] > A[\text{largest}]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



So largest = i

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i


if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



↑ ↑ ↑
i=0 l=1 r=2
↑

largest=0

$A[r] > A[\text{largest}]$

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.\text{size}$ and $A[l] > A[i]$

largest = l

else

largest = i


if $r \leq A.\text{size}$ and $A[r] > A[\text{largest}]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



↑ ↑ ↑
i=0 l=1 r=2
↑

largest=0

So update largest

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



largest isn't i

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if **largest** \neq i

swap A[i] and A[largest]

Heapify(A, largest)



So swap with i

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



And recurse

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



Find children

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



l is off the end of the heap

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



So largest is i

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



r is also off the end of the heap

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap $A[i]$ and $A[largest]$

Heapify(A, largest)



largest=i

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if **largest** \neq i

swap A[i] and A[largest]

Heapify(A, largest)



So we're done

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i



if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



next time:

applications of binary heaps