

<http://xkcd.com/835/>

Lecture 14

Applications of priority queues

EECS-214

Priority queues

- Like normal queues
 - Objects wait in line to be processed
- However, items have an associated **numeric priority**
 - Priority specified when added to queue
 - Objects removed from queue in order of priority
- Slightly different API
 - **Insert**(object, priority)
 - Adds object with specified priority
 - **ExtractMax**()
 - Returns highest priority object




Priority queues

- Like normal queues
 - Objects wait in line to be processed
- However, items have an associated numeric priority
 - Priority specified when added to queue
 - Objects removed from queue in order of priority
- Slightly different API
 - Insert(object, priority)
 - Adds object with specified priority
 - Extract**Min**()
 - “**Min priority queue**”
 - Returns **lowest** priority object
 - Except that in most applications, we usually want the lower “priorities” first




Heaps



- Heaps are a simple **tree structure** for implementing priority queues
- Rather than requiring their in-order traversal to be sorted
 - We just require that **parent nodes be larger than** their child nodes
- There are lots of exotic types of heaps
 - We'll focus on binary heaps
 - Which are **complete binary trees** with the heap property
 - We'll get to the completeness thing in a minute...

Heaps




Proposition: the largest element of a heap is always its root

Proof:

- Suppose some other element is the largest element
- Since it isn't the root, it must have a parent
- Since it's the largest element, it must be larger than its parent
- But that contradicts the definition of a heap
- So the largest element must be the root


Binary heaps

- A **binary heap** is a
 - **Complete** binary tree
 - That satisfies the heap property
- Great!
- How do we ensure that the heap is a complete binary tree?




Embedding in an array

- It turns out that any complete binary tree can be **embedded an array** in a particularly clever way
- We can compute
 - The position of its parent in the array,
 - and the positions of its children,
 - directly from its own position



Embedding in an array

- Store the **root** in the **first element** (element 0)
- For any node
 - Let i be its position in the array (for the root, $i = 0$)
 - Store its **left child** at position $2i + 1$
 - Store its **right child** at position $2i + 2$
 - Its **parent** can be found at position $\lfloor (i - 1)/2 \rfloor$
- Trust me that this works :-)



Heap insertion using the array representation

HeapInsert(A, value)

A.size = A.size + 1

i = A.size



while i > 0 and

A[Parent(i)] < key

A[i] = A[Parent(i)]

i = Parent(i)

A[i] = key



Extracting an element

HeapExtractMax(A, value)


max = A[0]

A[0] = A[A.size]

A.size--

Heapify(A,0)

return max



Extracting an element

Heapify(A, i)

l = Left(i)

r = Right(i)

if $l \leq A.size$ and $A[l] > A[i]$

largest = l

else

largest = i


if $r \leq A.size$ and $A[r] > A[largest]$

largest = r

if largest \neq i

swap A[i] and A[largest]

Heapify(A, largest)



Sorting

[Pretend I had found some awesome clipart about sorting here]

- Priority queues can be used for **sort algorithms**
 - **Add all items** to the queue
 - Repeatedly **extract max**
 - Or min, if it's a min queue
 - **Write them in order**

Version 1

- Here's the **straightforward** way of sorting using a heap
 - We **make the heap**
 - **Insert** all the **elements** into it
 - Pull them out one at a time, and **write them back** into A
- However, we're **copying** all the data
 - From one array, A
 - To the heap, H, which is also an array
 - **$O(n)$ space**

Heapsort(A)

H = new, empty heap

for each e in A

 Insert(H, e)

for i = A.Length-1 to 0

 A[i] = ExtractMax(H)

In-place heapsort

- It would be cooler if we could **build the heap inside A itself**
- How do we do that?

HeapSort(A)

H = new, empty heap

for each e in A

Insert(H, e)

for i = A.Length-1 to 0

A[i] = ExtractMax(H)

Sketch of in-place heap construction



Sketch of in-place heap construction

- Start with the array

4	6	2	7	3	1	5
---	---	---	---	---	---	---



Sketch of in-place heap construction

- Start with the array
- Pretend that it's a binary tree
- It probably **doesn't satisfy** the heap property
 - i.e. there are probably nodes that are larger than their parents





Sketch of in-place heap construction

- But we can think of **each leaf** as a little **1-element heap**
 - And they **automatically satisfy** the heap property
 - Because they only have one element





Sketch of in-place heap construction

- Now run **heapify** on each of their parents
- Heapify
 - **Checks** if the **parent** is larger than both children





Sketch of in-place heap construction

- Now run **heapify** on each of their parents
- Heapify
 - Checks if the **parent** is larger than both children
 - If not, it **swaps** it with the **larger child**





Sketch of in-place heap construction

- Now we have
 - A bunch of **2-level subtrees** that
 - Each **satisfy the heap property**





Sketch of in-place heap construction

- **Repeat** at the next level





Sketch of in-place heap construction

- **Repeat** at the next level
- This may require Heapify to **recurse**





Sketch of in-place heap construction

- Repeat at the next level
- This may require Heapify to recurse
- But when we've **done every level**



Sketch of in-place heap construction

- Repeat at the next level
- This may require Heapify to recurse
- But when we've **done every level**
- We've transformed the data, **in-place**, into a **binary heap**



Algorithm for in-place heap construction

- The **last half** of the array is guaranteed to be **leaves**
 - So we don't have to do anything with it

```
BuildHeap(A)
  for i = A.Length/2 to 1
    Heapify(A, i, A.Length)
```

Algorithm for in-place heap construction

- We want to call **Heapify**
 - On every **non-leaf** node
 - **Starting at the bottom** of the tree
 - And moving upwards
- Since **lower parts** of the tree are **at the end** of the array,
 - All we have to do is **start halfway** through the array
 - and **move back**
 - Calling Heapify

BuildHeap(A)

for $i = A.Length/2$ to 0

 Heapify(A, i, A.Length)

Version 2

Now this is really kind of **cool**

- We **build** the heap **in place**
- Then we **loop**
 - Each iteration **removes** the **maximal** element
 - That **shrinks** the heap
 - **Making room** at the end of the array
 - Which is **where** we would want to **put** the maximal element, anyway!

Heapsort(A)

 BuildHeap(A)


 for i = A.Length-1 to 1

 A[i] = ExtractMax(H)

Completing the sort

ExtractMax



- Grabs the maximal element, 7



Completing the sort

ExtractMax



- Grabs the maximal element, 7
- Replaces it with the last leaf, 2



Completing the sort

ExtractMax



- Grabs the maximal element, 7
- Replaces it with the last leaf, 2
- And calls Heapify



Completing the sort

ExtractMax



- Grabs the maximal element, 7
- Replaces it with the last leaf, 2
- And calls Heapify



Completing the sort

ExtractMax



- Grabs the maximal element, 7
- Replaces it with the last leaf, 2
- And calls Heapify



Completing the sort

ExtractMax



- Grabs the maximal element, 7
- Replaces it with the last leaf, 2
- And calls Heapify



Completing the sort

ExtractMax



- Grabs the maximal element, 7
- Replaces it with the last leaf, 2
- And calls Heapify
 - Thereby reestablishing the heap property



Completing the sort



ExtractMax

- Grabs the maximal element, 7
- Replaces it with the last leaf, 2
- And calls Heapify
 - Thereby reestablishing the heap property
- And now we store 7 (the max) at the end of the array)





Completing the sort

Repeat





Completing the sort

- Grab the max, 6





Completing the sort

- Grab the max, 6
- Replace it with the last leaf, 1





Completing the sort

- Grab the max, 6
- Replace it with the last leaf, 1





Completing the sort

- Grab the max, 6
- Replace it with the last leaf, 1
- Heapify





Completing the sort

- Grab the max, 6
- Replace it with the last leaf, 1
- Heapify





Completing the sort

- Grab the max, 6
- Replace it with the last leaf, 1
- Heapify





Completing the sort

- Grab the max, 6
- Replace it with the last leaf, 1
- Heapify
- Store the max, 6, in the next-to-last slot





Completing the sort

Repeat





Completing the sort

- Grab the max, 5





Completing the sort

- Grab the max, 5
- Replace with the last leaf, 3





Completing the sort

- Grab the max, 5
- Replace with the last leaf, 3





Completing the sort

- Grab the max, 5
- Replace with the last leaf, 3
- Heapify





Completing the sort

- Grab the max, 5
- Replace with the last leaf, 3
- Heapify





Completing the sort

- Grab the max, 5
- Replace with the last leaf, 3
- Heapify





Completing the sort

- Grab the max, 5
- Replace with the last leaf, 3
- Heapify
- Store the max





Completing the sort

Repeat





Completing the sort

- Grab the max, 4





Completing the sort

- Grab the max, 4
- Replace with the last leaf, 2





Completing the sort

- Grab the max, 4
- Replace with the last leaf, 2





Completing the sort

- Grab the max, 4
- Replace with the last leaf, 2
- Heapify





Completing the sort

- Grab the max, 4
- Replace with the last leaf, 2
- Heapify





Completing the sort

- Grab the max, 4
- Replace with the last leaf, 2
- Heapify




Completing the sort

- Grab the max, 4
- Replace with the last leaf, 2
- Heapify
- Store the max




Completing the sort

Repeat




Completing the sort

- Grab the max, 3




Completing the sort

- Grab the max, 3
- Replace with the last leaf, 1



Completing the sort


- Grab the max, 3
- Replace with the last leaf, 1



1	2	1	4	5	6	7
---	---	---	---	---	---	---

Completing the sort


- Grab the max, 3
- Replace with the last leaf, 1
- Heapify



1	2	1	4	5	6	7
---	---	---	---	---	---	---


Completing the sort

- Grab the max, 3
- Replace with the last leaf, 1
- Heapify




Completing the sort

- Grab the max, 3
- Replace with the last leaf, 1
- Heapify
- Store the max




Completing the sort

Repeat




Completing the sort

- Grab the max, 2



Completing the sort

- Grab the max, 2
- Replace with the last leaf, 1



Completing the sort

- Grab the max, 2
- Replace with the last leaf, 1



Completing the sort

- Grab the max, 2
- Replace with the last leaf, 1
- Heapify

1

1	1	3	4	5	6	7
---	---	---	---	---	---	---

Completing the sort

- Grab the max, 2
- Replace with the last leaf, 1
- Heapify

1

1	1	3	4	5	6	7
---	---	---	---	---	---	---

Completing the sort

- Grab the max, 2
- Replace with the last leaf, 1
- Heapify
- Store the max

1

1	2	3	4	5	6	7
---	---	---	---	---	---	---

Completing the sort

- Only **one element left**
- So it has to be the **smallest** element
- And it has to be in **position 0**

1

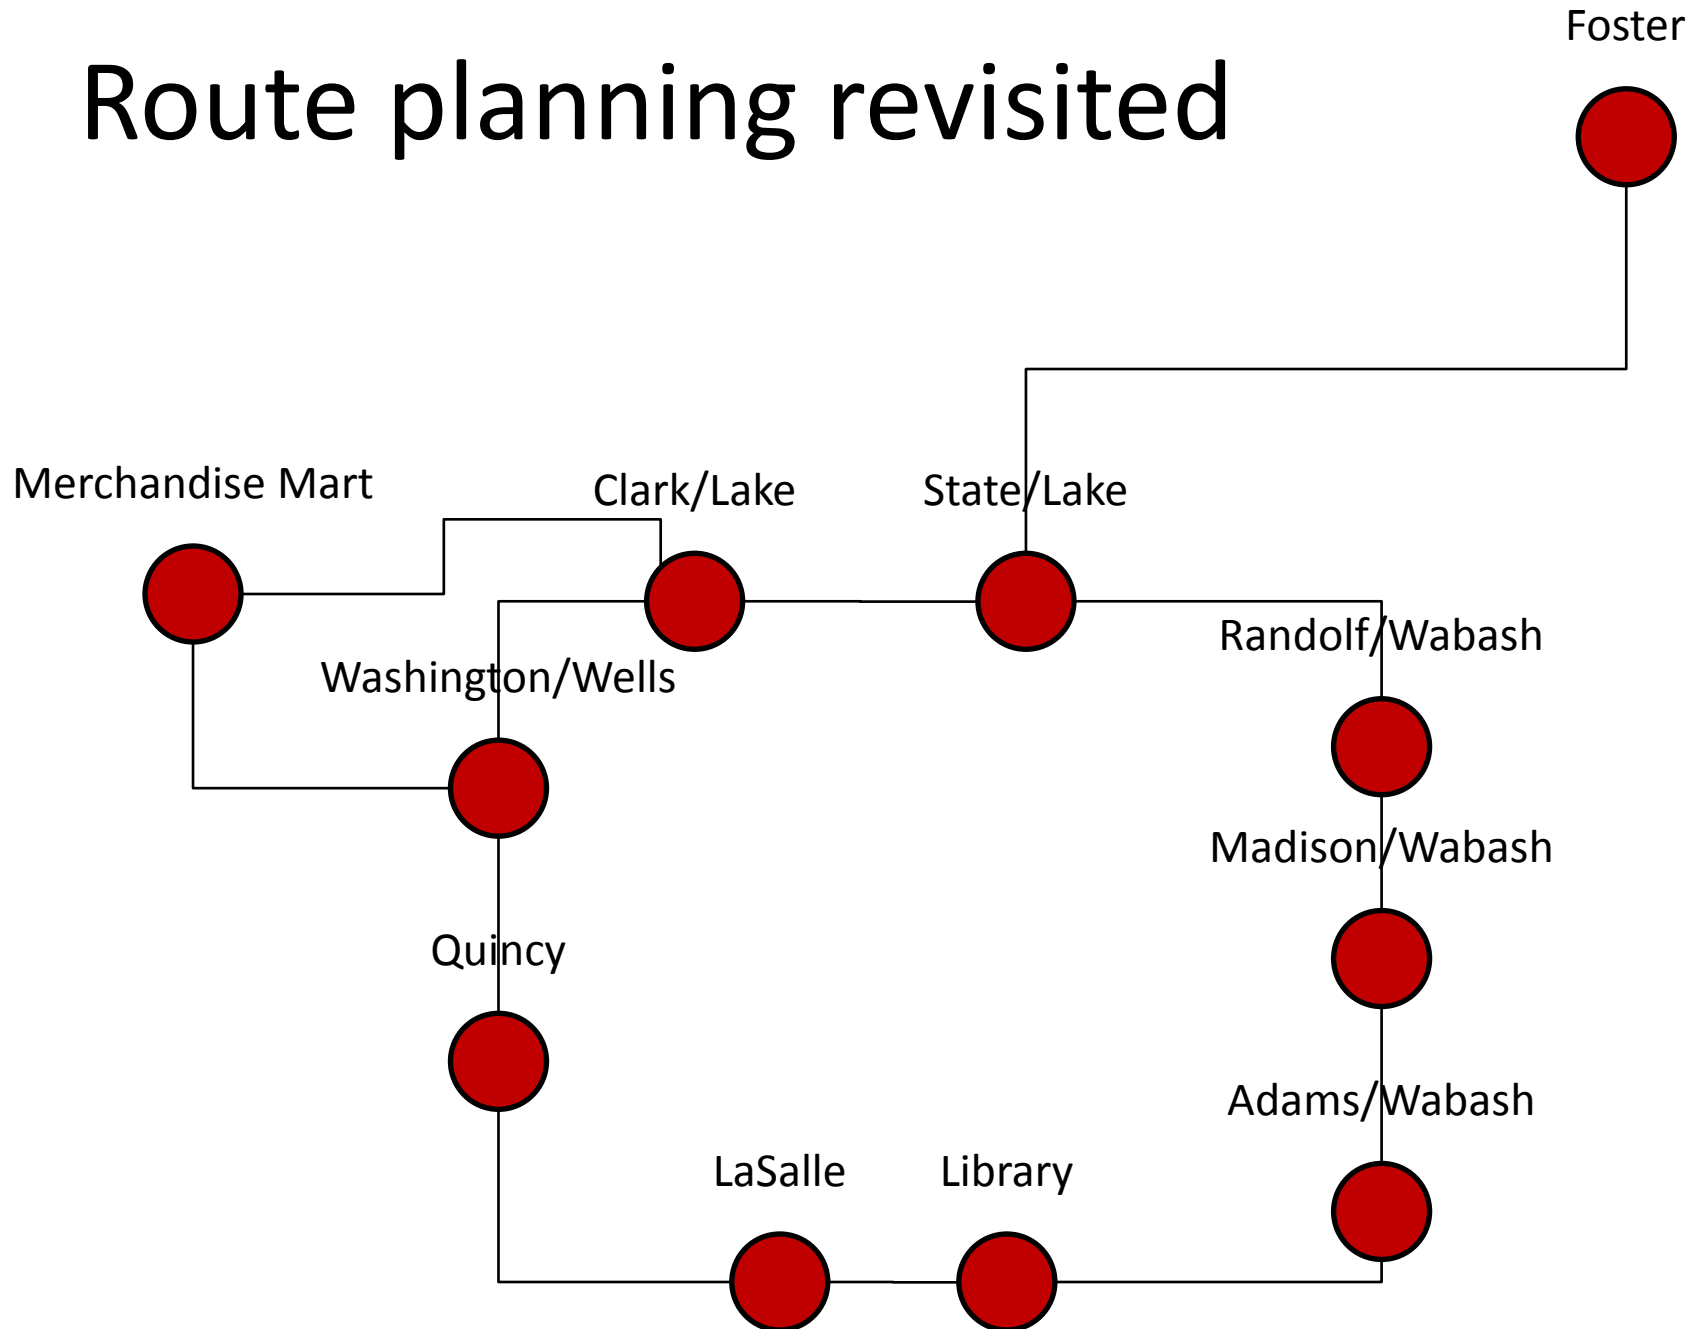
We're done!

1	2	3	4	5	6	7
---	---	---	---	---	---	---

Analysis

- Every **heap operation** is $O(\log n)$
- We do $n/2$ calls to Heapify to build the heap
- We do n calls to ExtractMax to do the final sorting
- So execution time is $O(n \log n)$
- The nice thing about Heapsort is that it's $O(n \log n)$ **worst-case**, so it's **never quadratic**

Route planning revisited




Route planning revisited

- We used **breadth-first** search before to find a **shortest path** in a graph

- But BFS treats each edge as being equal

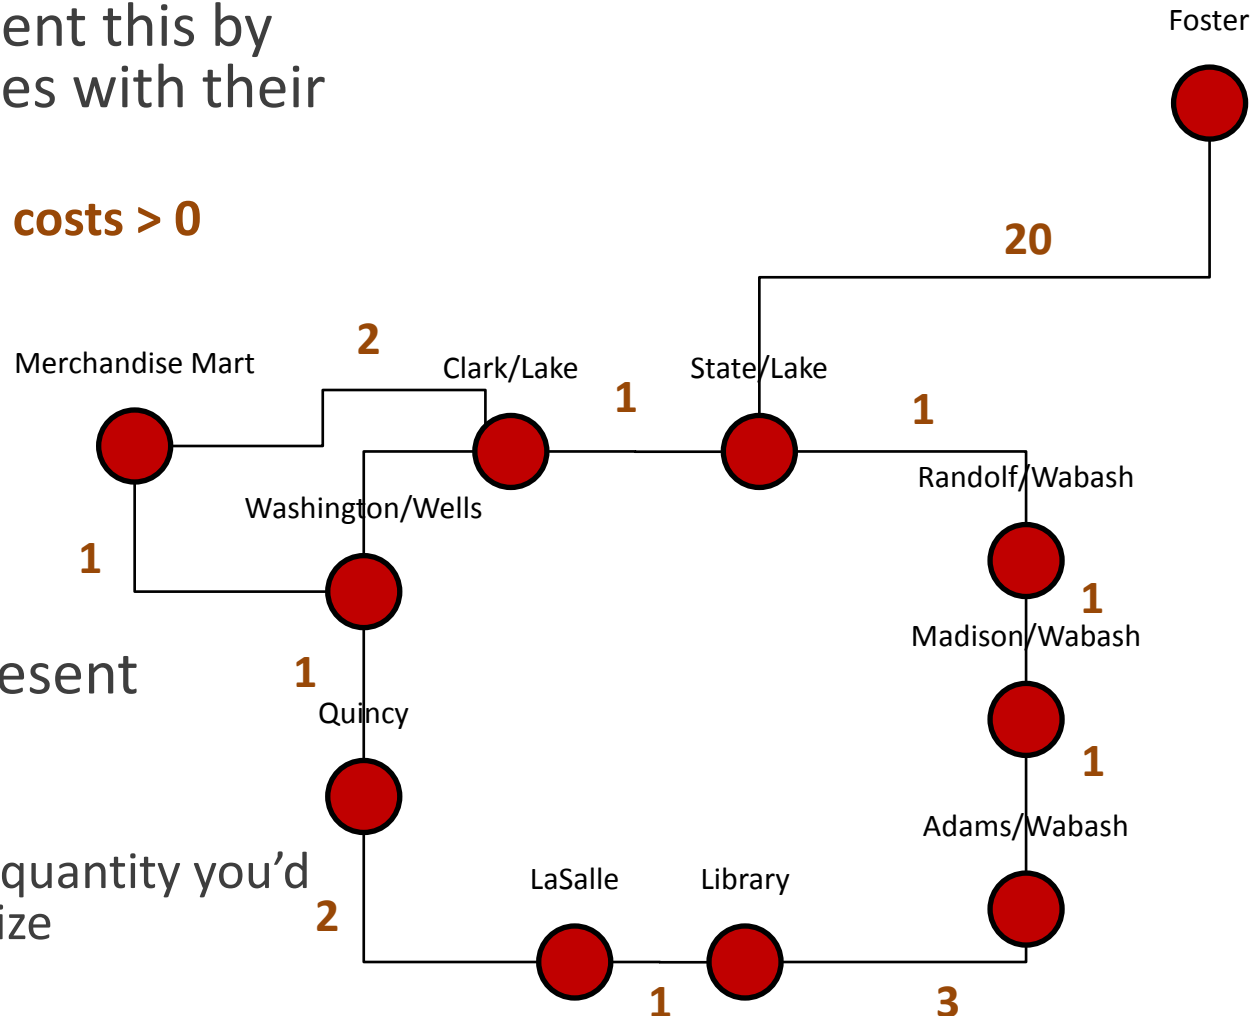
- In route planning, different edges take **different amounts of time**



Route planning revisited

- We can represent this by **weighting** edges with their “**costs**”

- We’ll assume **costs** > 0



- Costs can represent
 - Time
 - Distance
 - Or any other quantity you’d like to minimize

Path finding with edge costs

- How do we make a version of BFS that searches nodes
 - In order of **increasing total path cost**
 - Rather than increasing **number of edges**
- Use a **priority queue!**
 - **Priority** of a node = **cost of path from start to that node**

Dijkstra's least-cost path algorithm

Dijkstra(G, s, e)

PQ = new priority queue

Set all node costs to infinity

s.cost = 0

for each node n in G

 PQ.Insert(n, n.cost)

while PQ not empty

 u = PQ.ExtractMin()

 if u == end then done!

 for each neighbor v of u

 w = weight of edge from u to v

 newCost = u.cost + w

 if newCost < v.cost

 PQ.DecreaseKey(v, newCost)

 v.cost = newCost

 v.predecessor = u

Wait a minute... decrease key?

- We have to add a new operation to our priority queue: **decreasing the priority** of an item already in the queue
 - Note that since we're using extract min, rather than extract max, decreasing the “priority” actually **moves it ahead** in the queue
- How could we implement decrease key?

Implementing DecreaseKey

- One way we could do it would be to:
 - **Remove** it (somehow)
 - **Reinsert** it
- But the insert algorithm
 - Adds it at the bottom of the heap
 - **Swaps it upward** until its priority is lower than its parent
 - At least for a min heap

```
HeapInsert(A, key)
    A.size = A.size + 1
    i = A.size
    while i > 0 and
        A[Parent(i)] > key
        A[i] = A[Parent(i)]
        i = Parent(i)
    A[i] = key
```

Implementing DecreaseKey

So DecreaseKey is actually **easy**:

- Just **move** the node **up**
- Until its in the **right place**
- Just **copy the code for insert**
 - And **remove** the part that starts by inserting it at the end

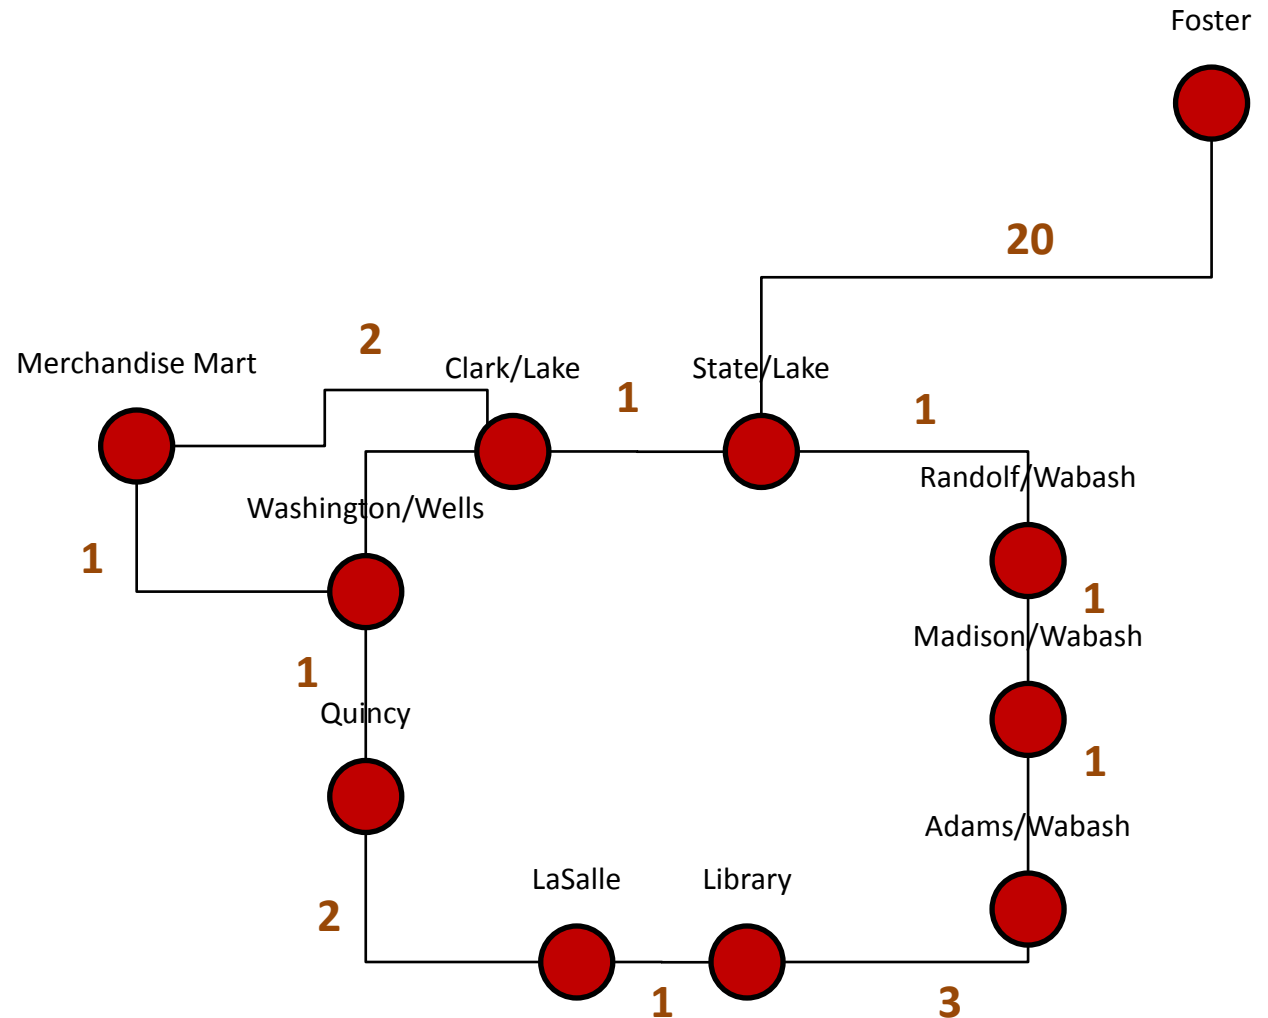
```
DecreaseKey(A, i, key)
    while i>0 and
        A[Parent(i)] > key
        A[i] = A[Parent(i)]
        i = Parent(i)
    A[i] = key
```

Implementing DecreaseKey

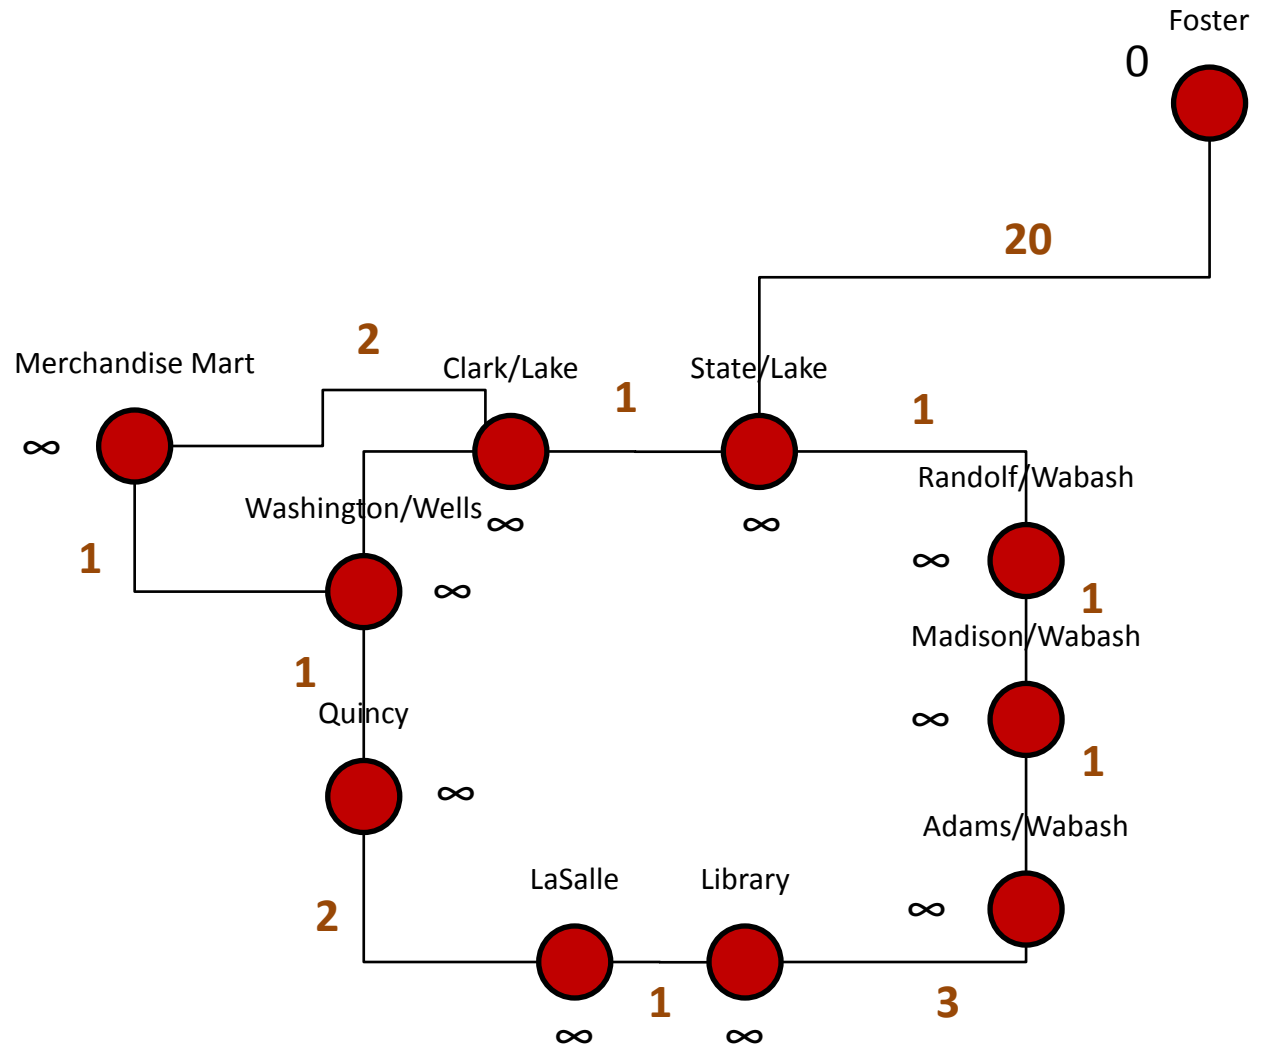
- Unfortunately, you need to know **where the node is** stored in the heap
 - i.e. its **index** in the stupid heap **array**
- Best done by storing the **index in the graph node** itself
- Have to remember to **update it** any time the node moves in the heap

```
DecreaseKey(A, i, key)
    while i > 0 and
        A[Parent(i)] > key
        A[i] = A[Parent(i)]
        i = Parent(i)
    A[i] = key
```

Running Dijkstra's algorithm



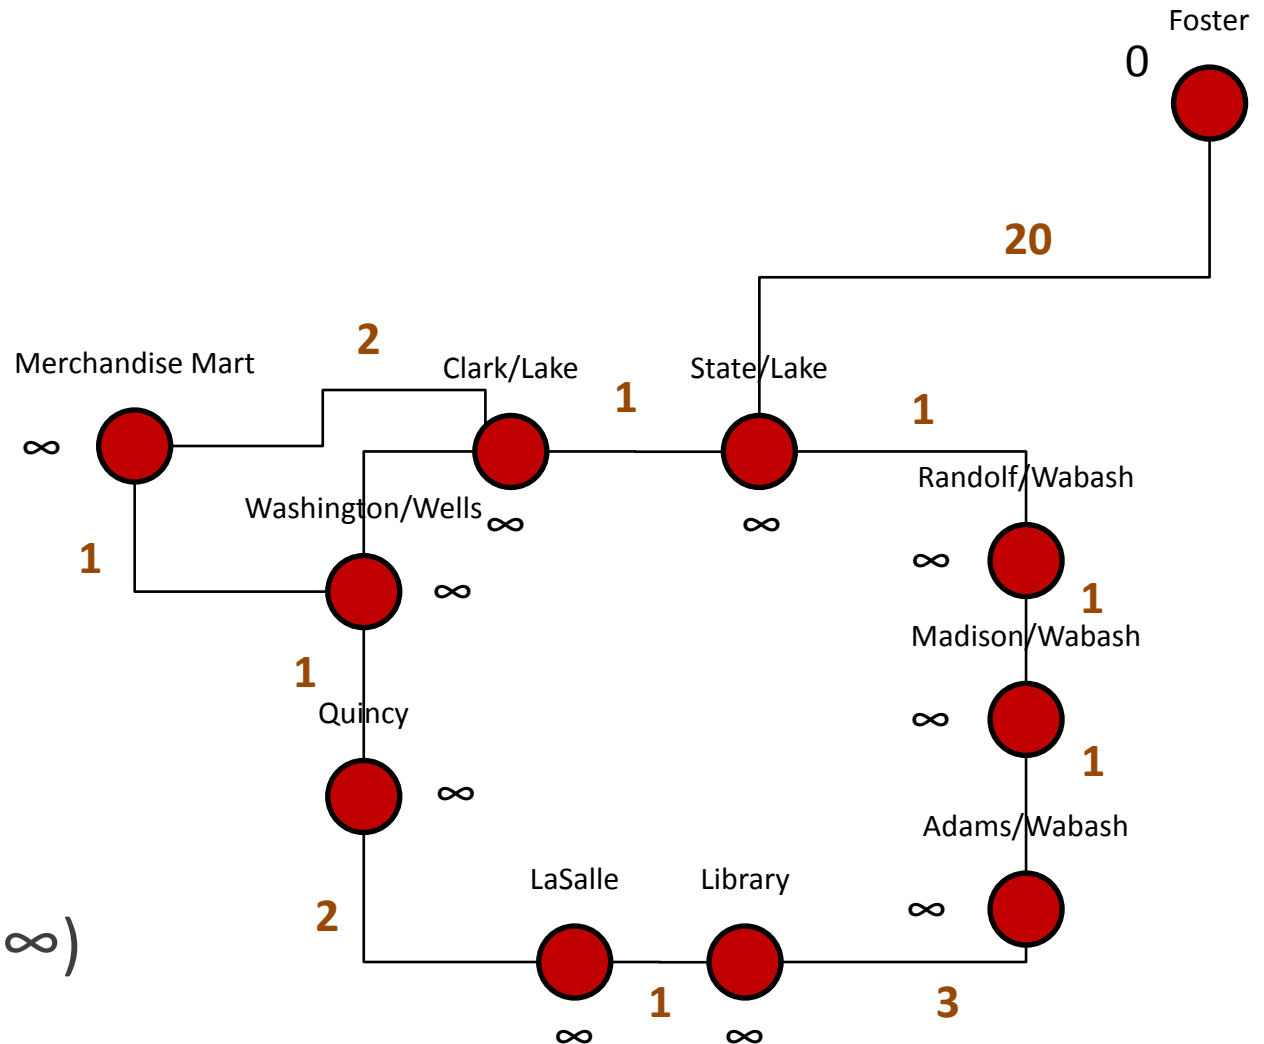
Initialize node costs



Initialize priority queue

Queue:

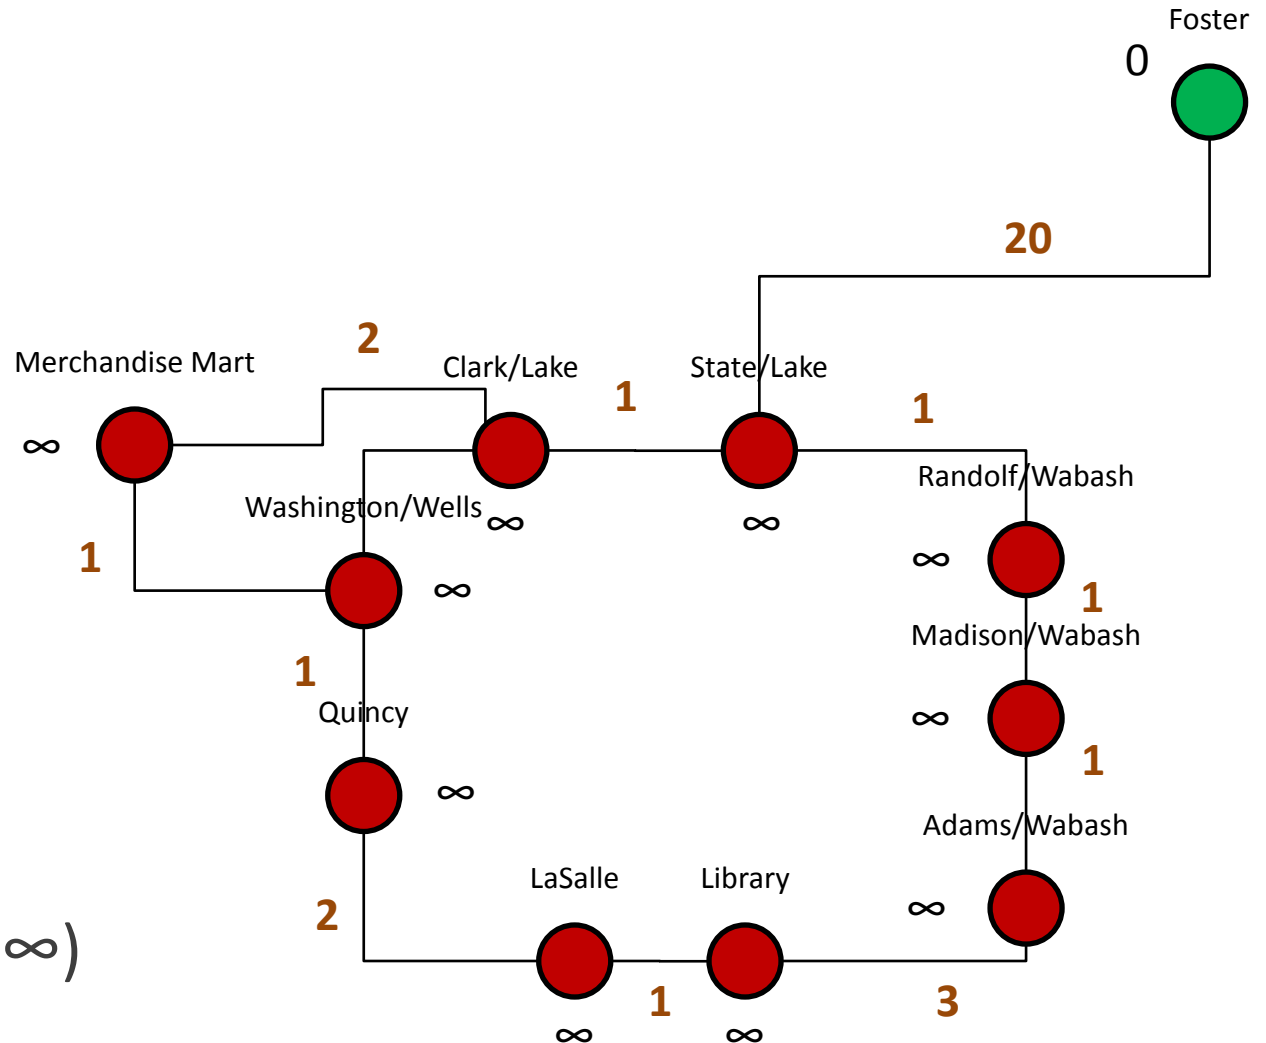
- Foster (0)
- Adams (∞)
- Clark (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Mmart (∞)
- Quincy (∞)
- Randolph (∞)
- State (∞)
- Washington (∞)



Extract min (Foster)

Queue:

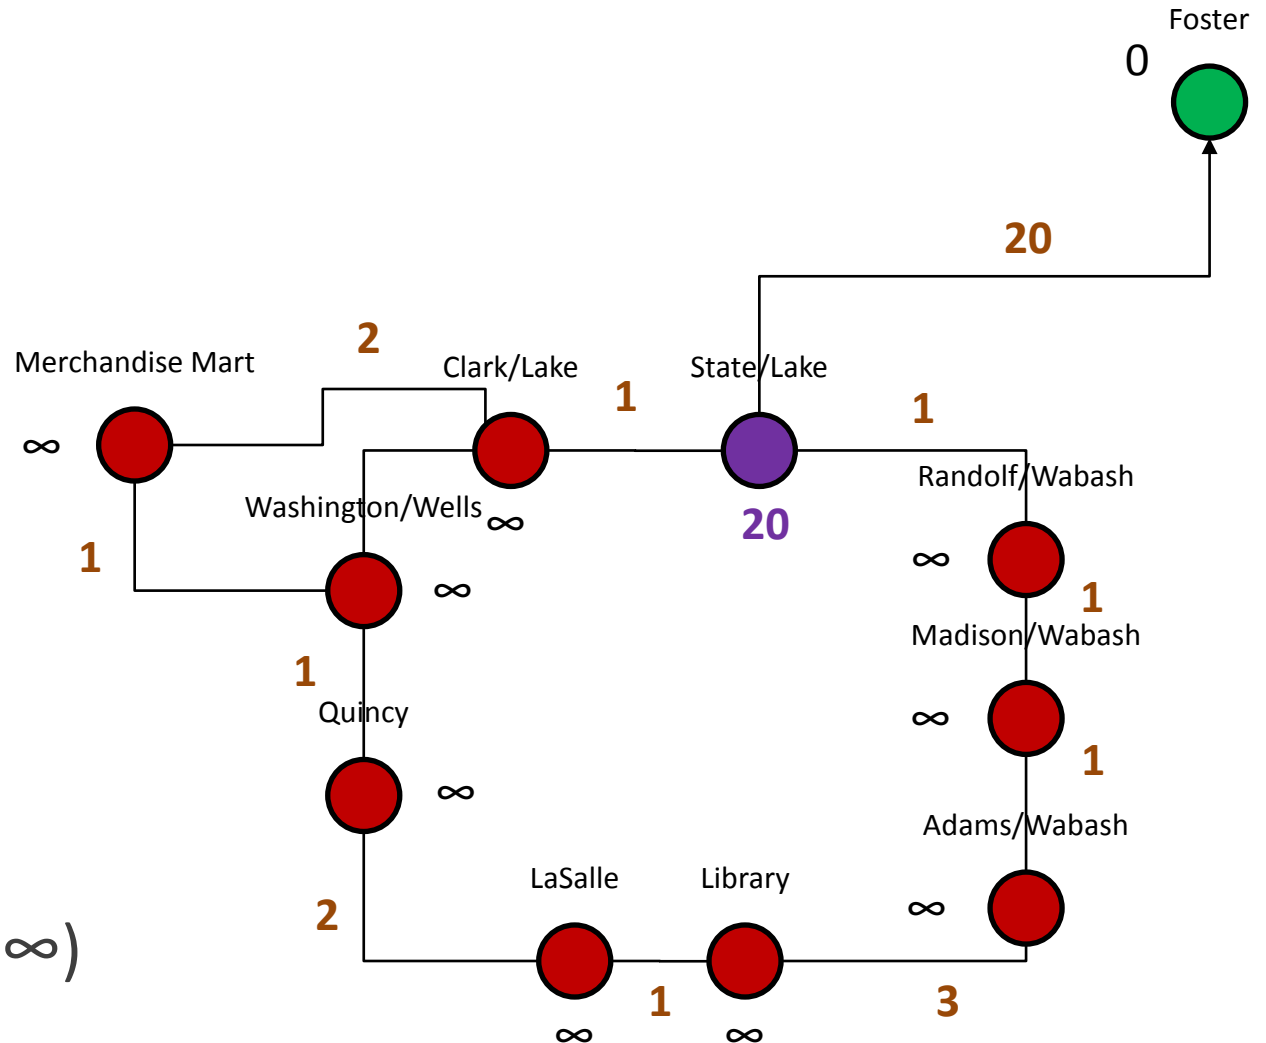
- Adams (∞)
- Clark (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Mmart (∞)
- Quincy (∞)
- Randolph (∞)
- State (∞)
- Washington (∞)



Update neighbors

Queue:

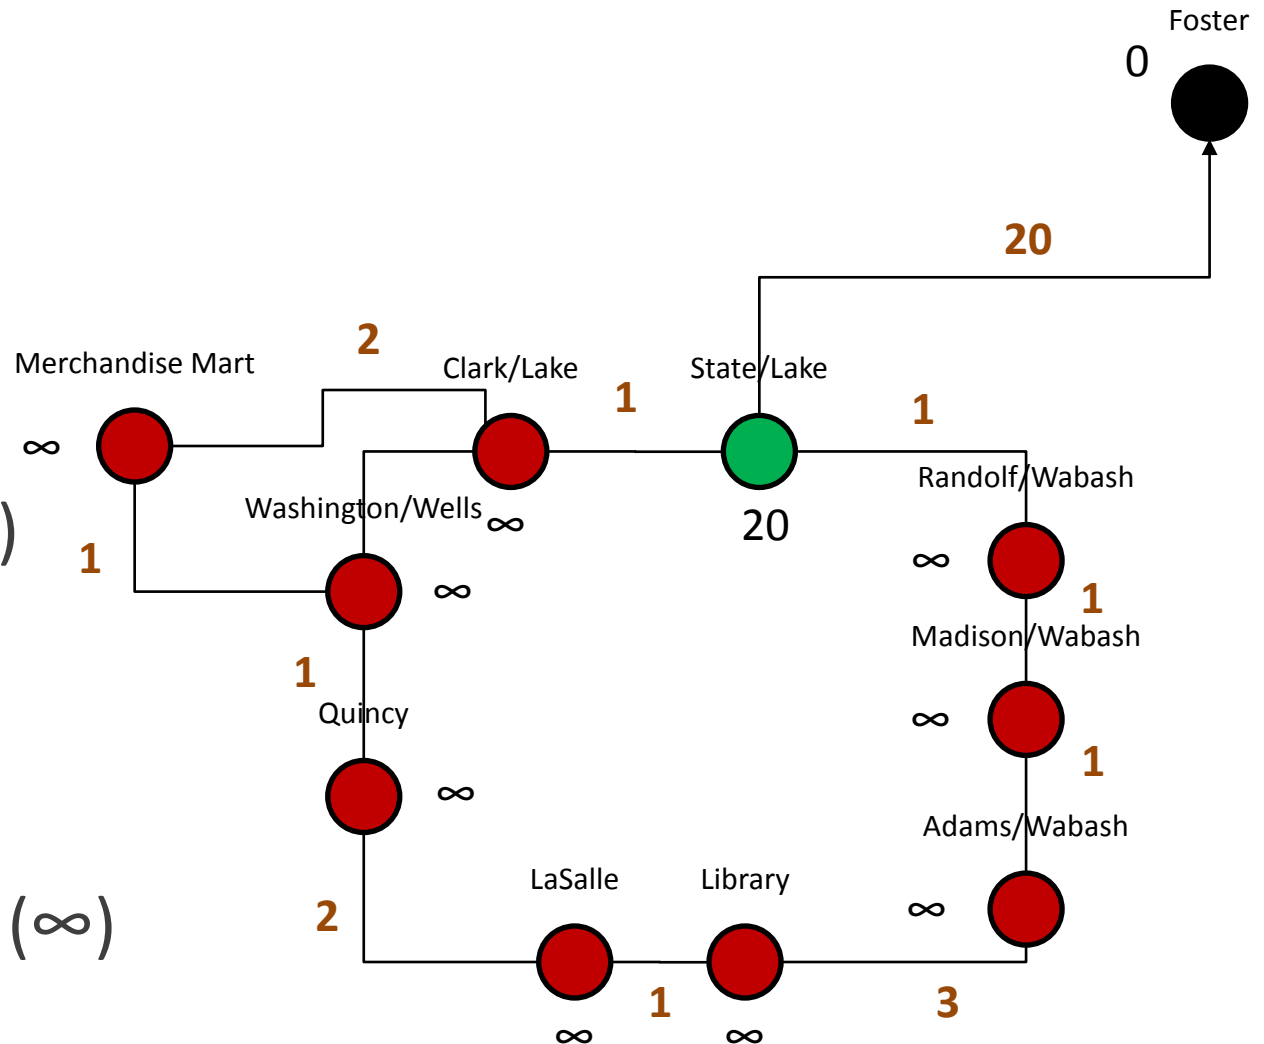
- **State (20)**
- Adams (∞)
- Clark (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Mmart (∞)
- Quincy (∞)
- Randolph (∞)
- Washington (∞)



Extract min (State)

Queue:

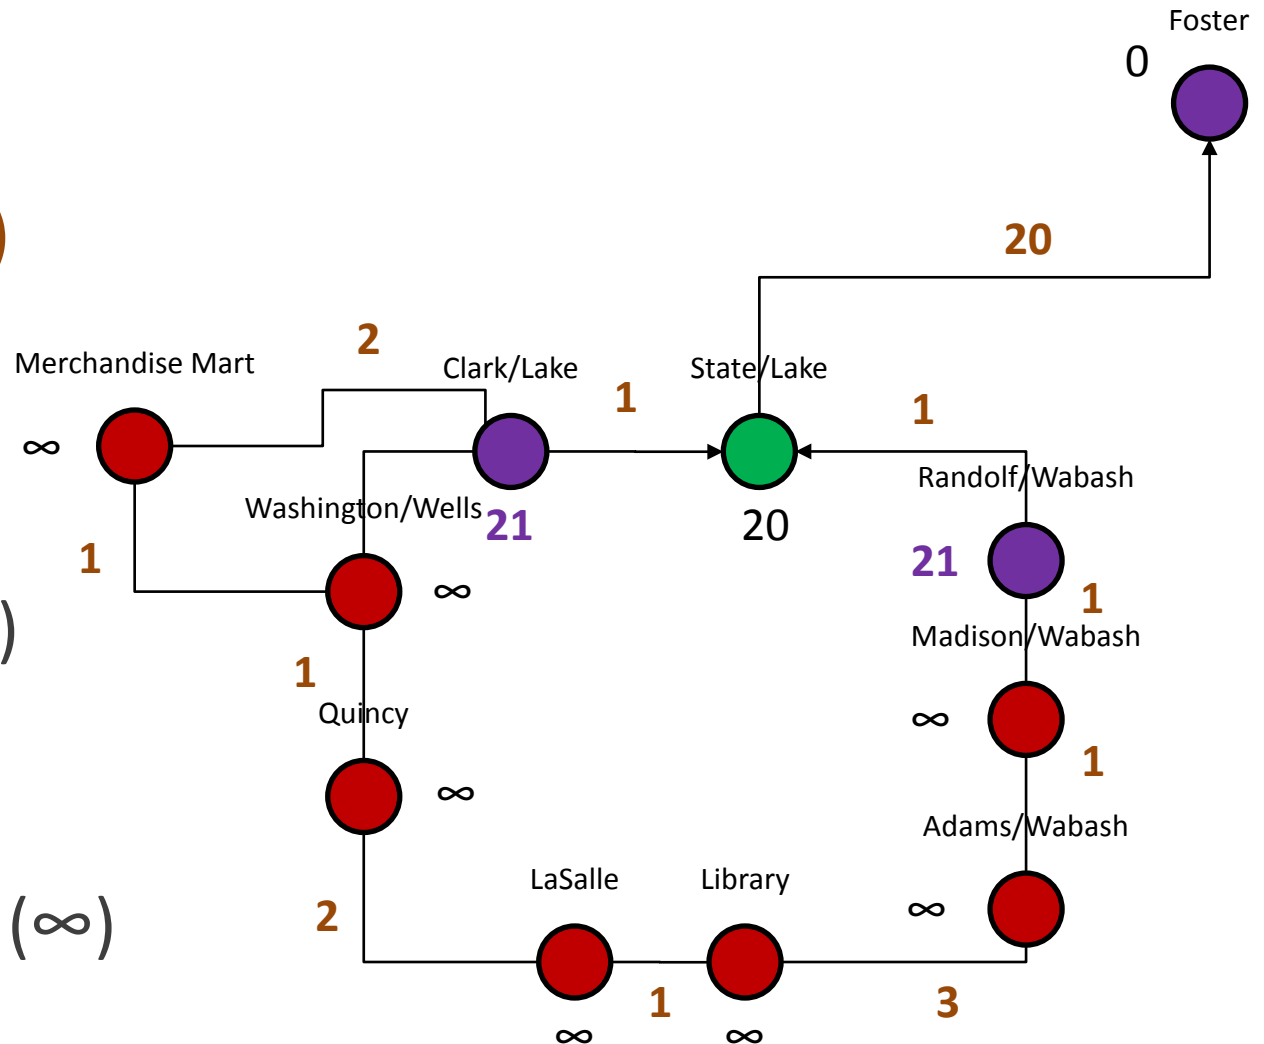
- Adams (∞)
- Clark (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Mmart (∞)
- Quincy (∞)
- Randolph (∞)
- Washington (∞)



Update neighbors

Queue:

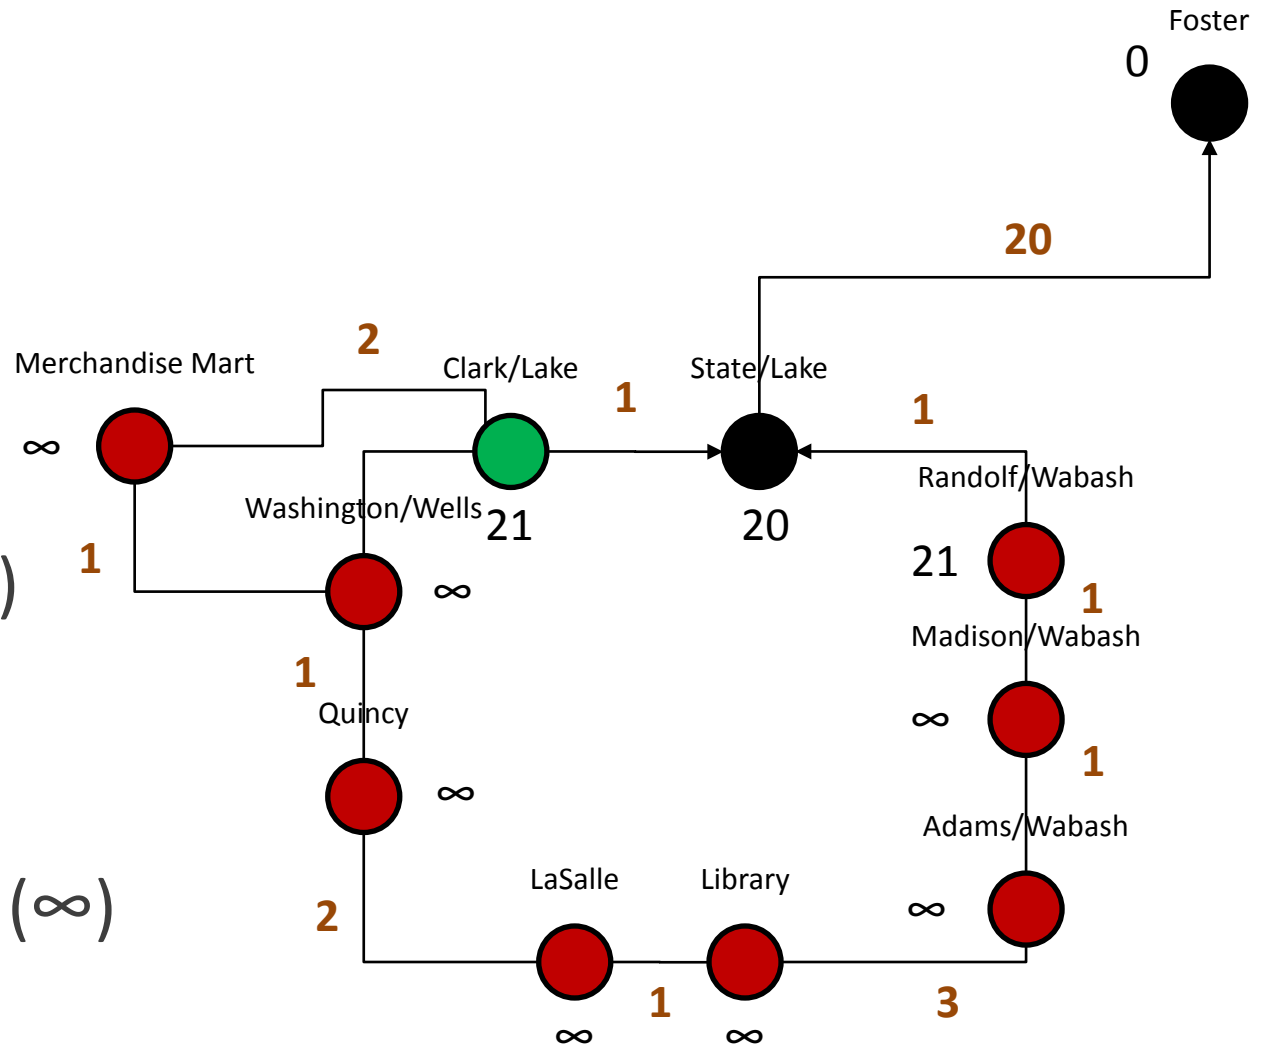
- **Clark (21)**
- **Randolf (21)**
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Mmart (∞)
- Quincy (∞)
- Washington (∞)



Extract min (Clark)

Queue:

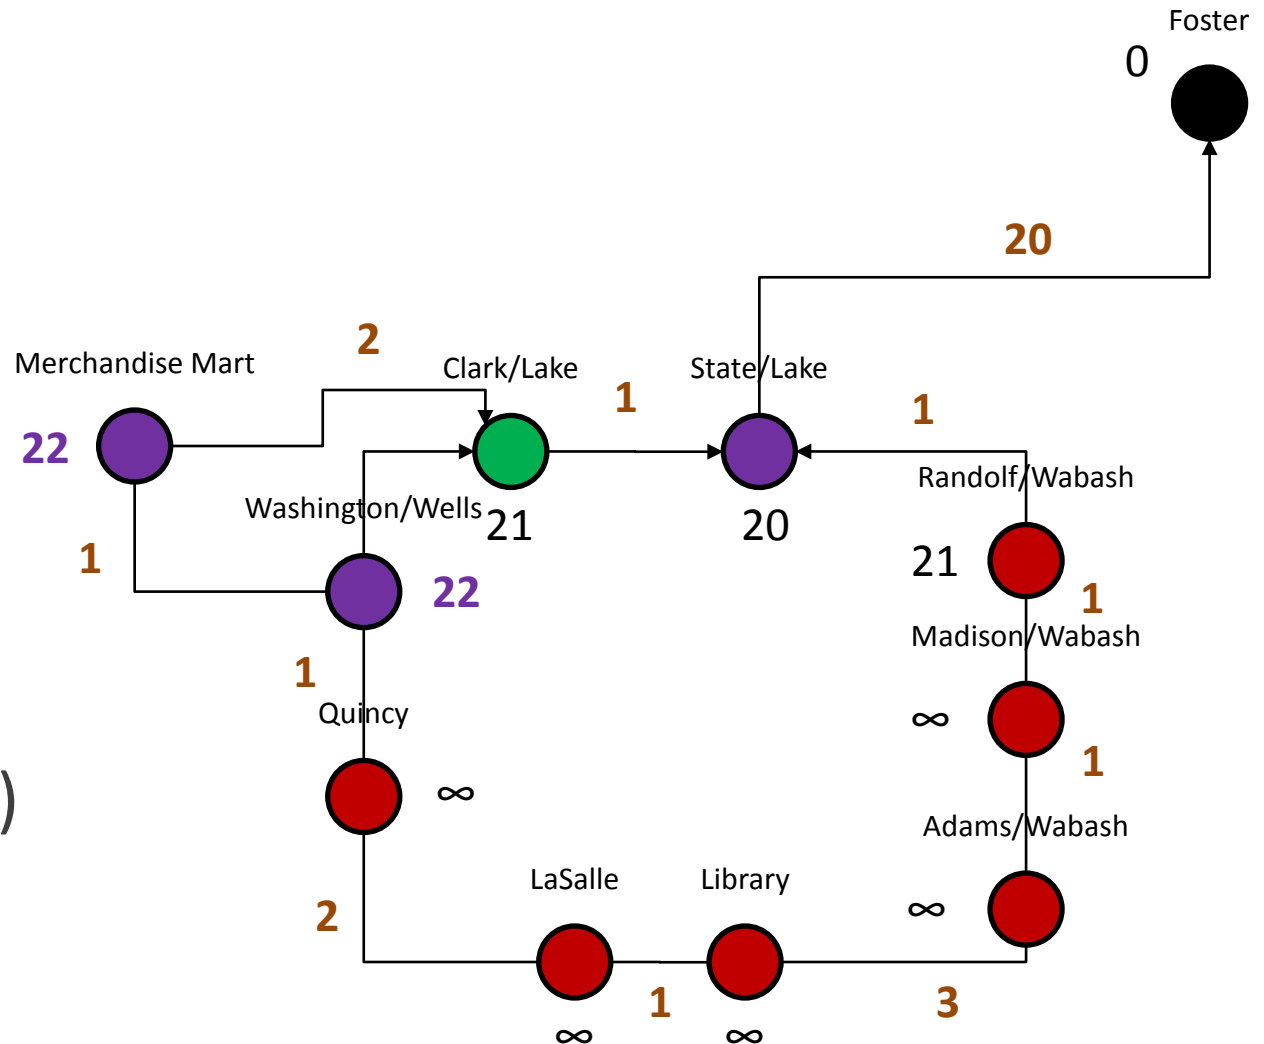
- Randolph (21)
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Mmart (∞)
- Quincy (∞)
- Washington (∞)



Update neighbors

Queue:

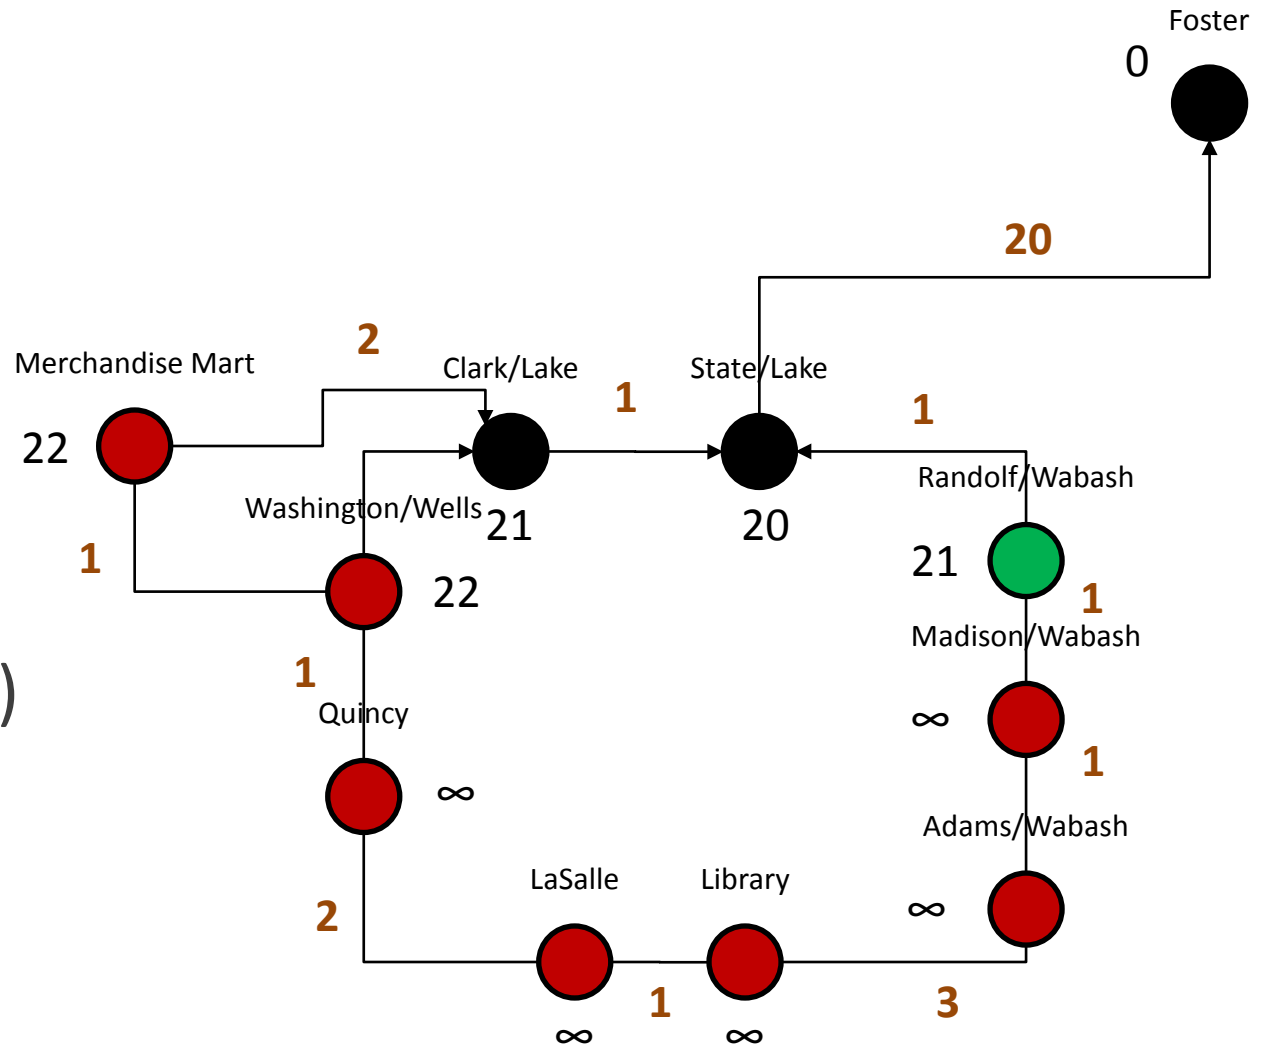
- Randolph (21)
- **Mmart (22)**
- **Wash (22)**
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Quincy (∞)



Extract min (Randolf)

Queue:

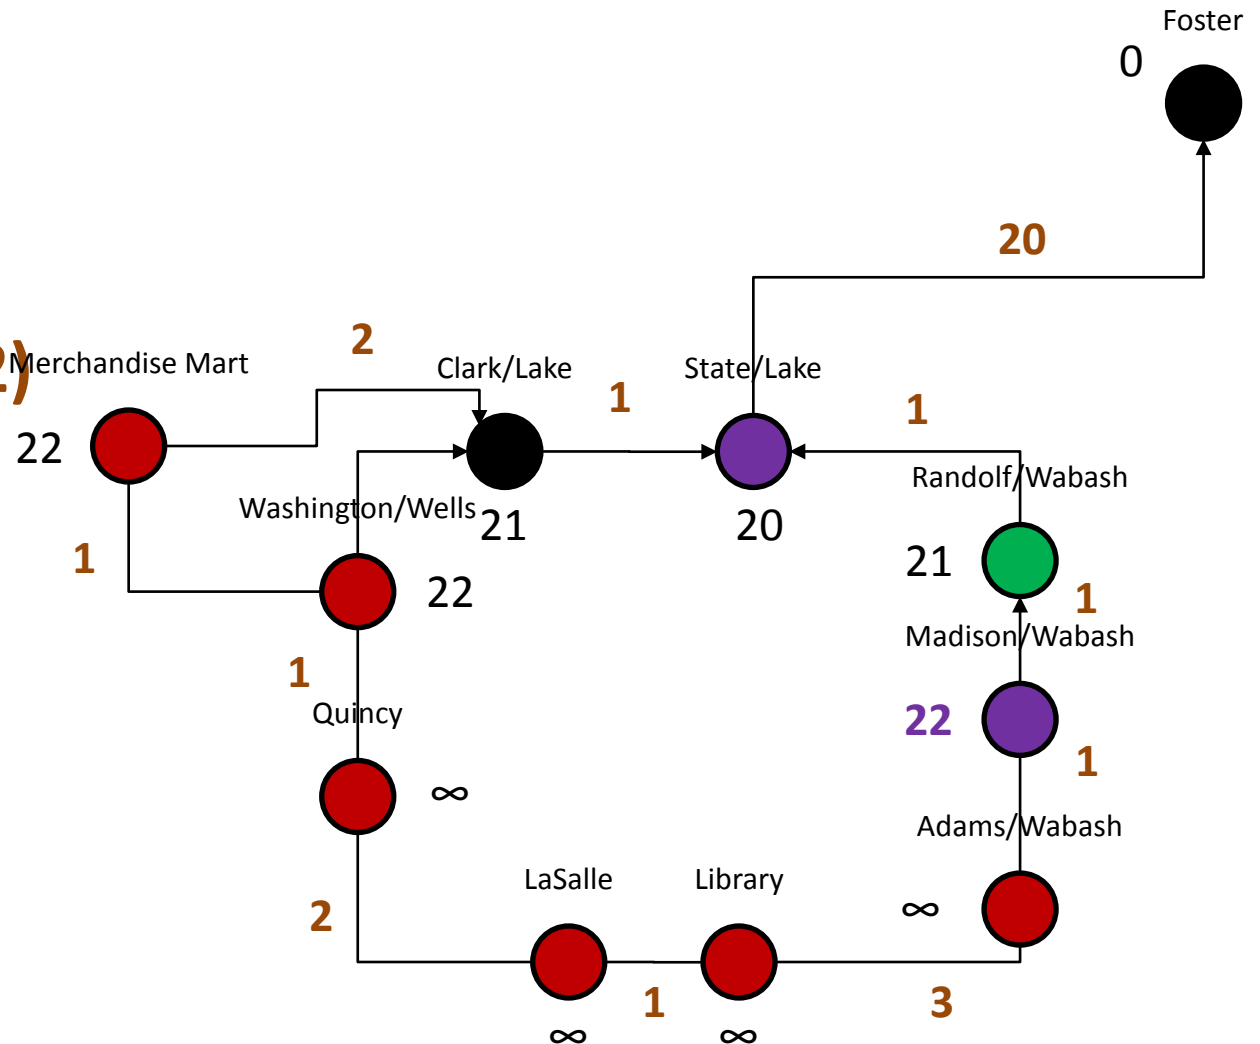
- Mmart (22)
- Wash (22)
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Madison (∞)
- Quincy (∞)



Update neighbors

Queue:

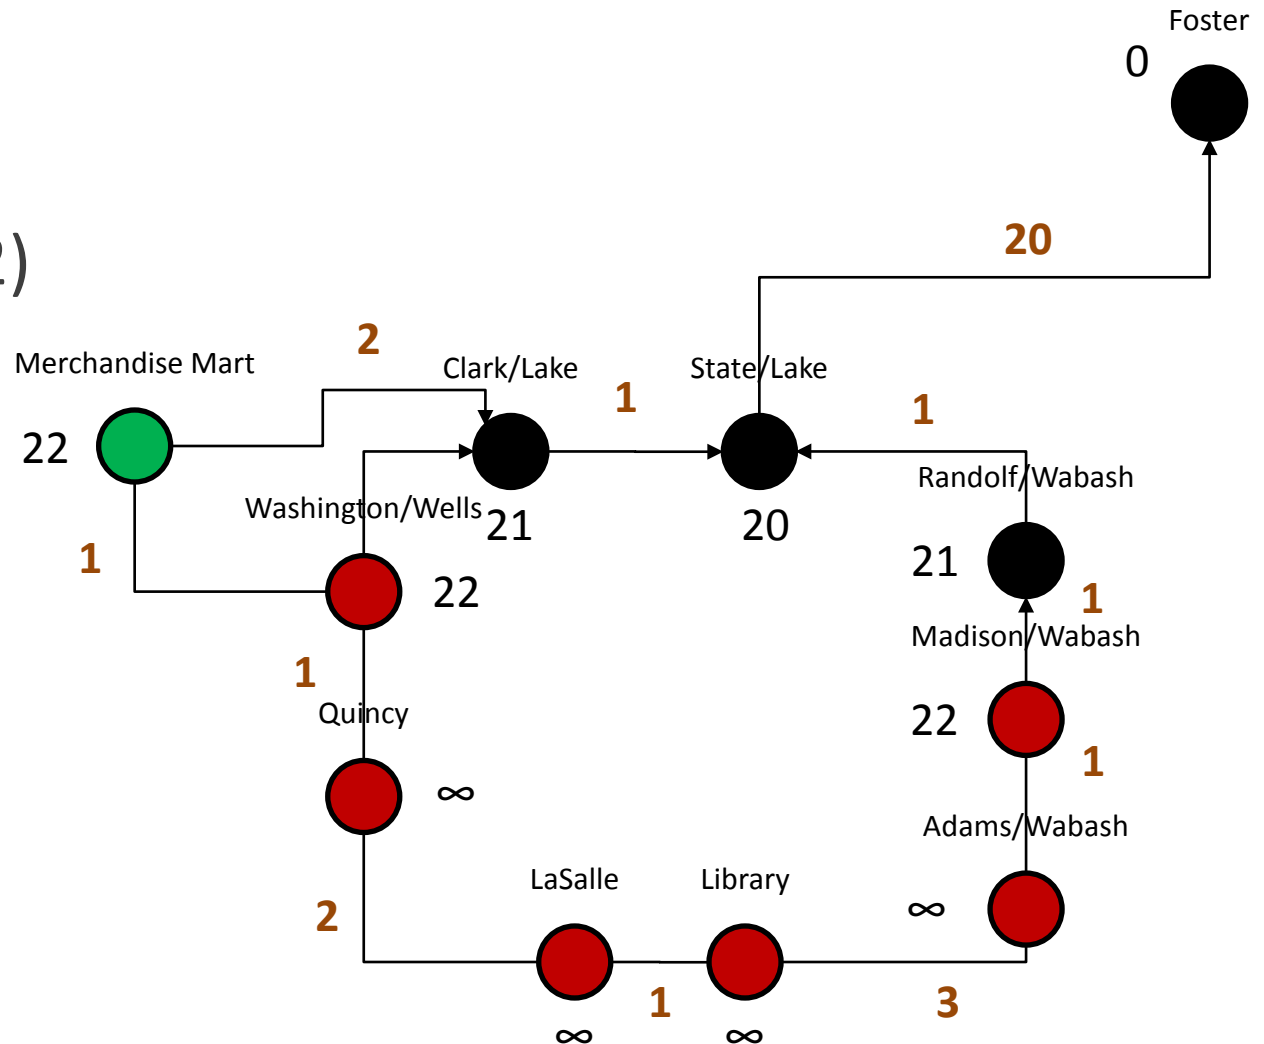
- Mmart (22)
- Wash (22)
- **Madison (22)**
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Quincy (∞)



Extract min (MMart)

Queue:

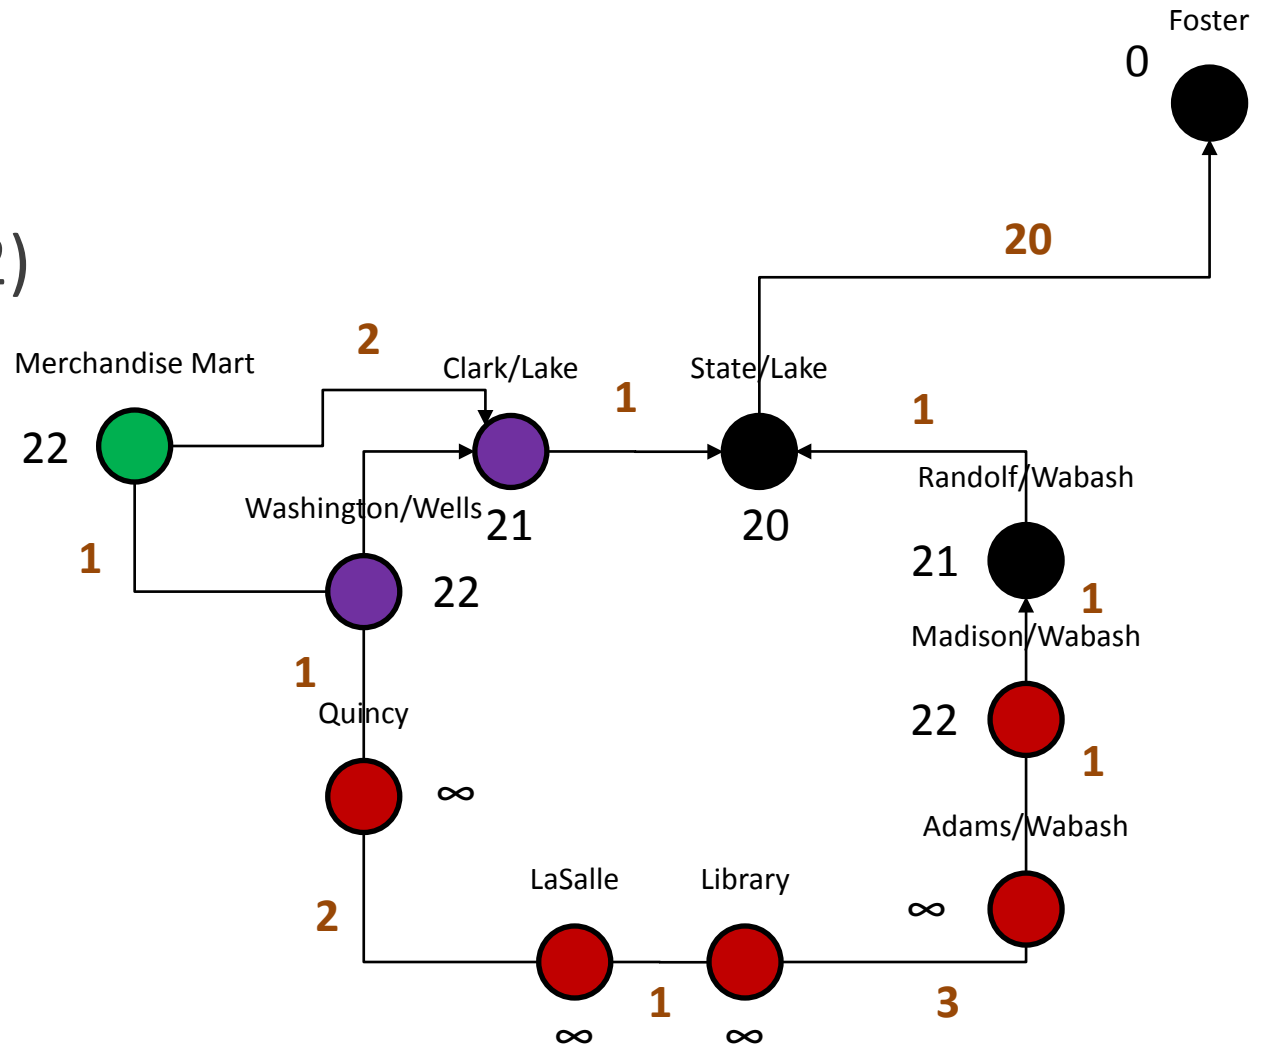
- Wash (22)
- Madison (22)
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Quincy (∞)



Update neighbors (but no distances changed)

Queue:

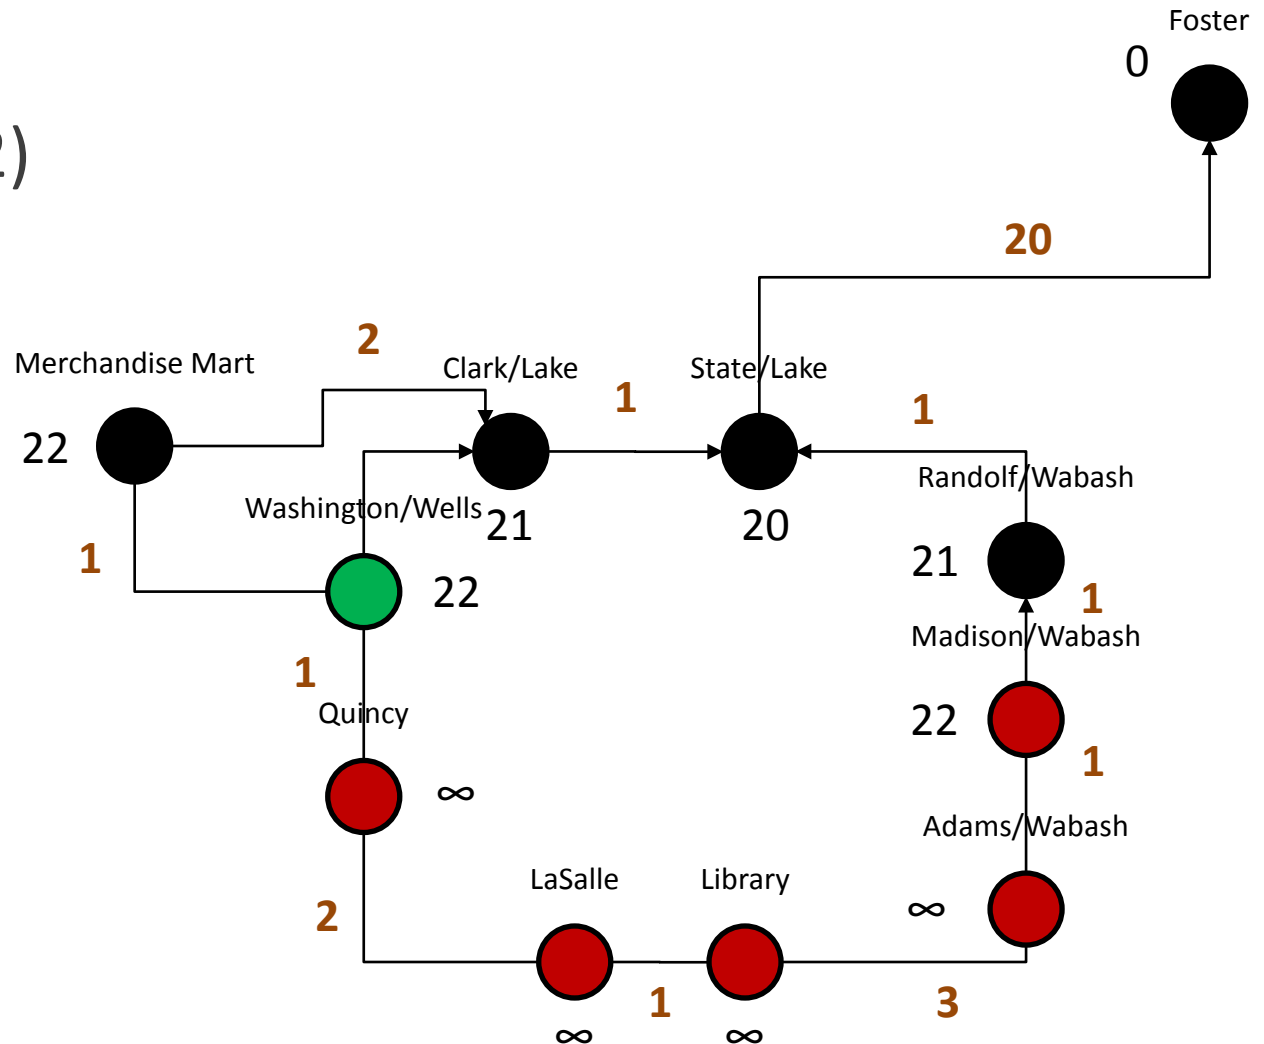
- Wash (22)
- Madison (22)
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Quincy (∞)



Extract min (Washington)

Queue:

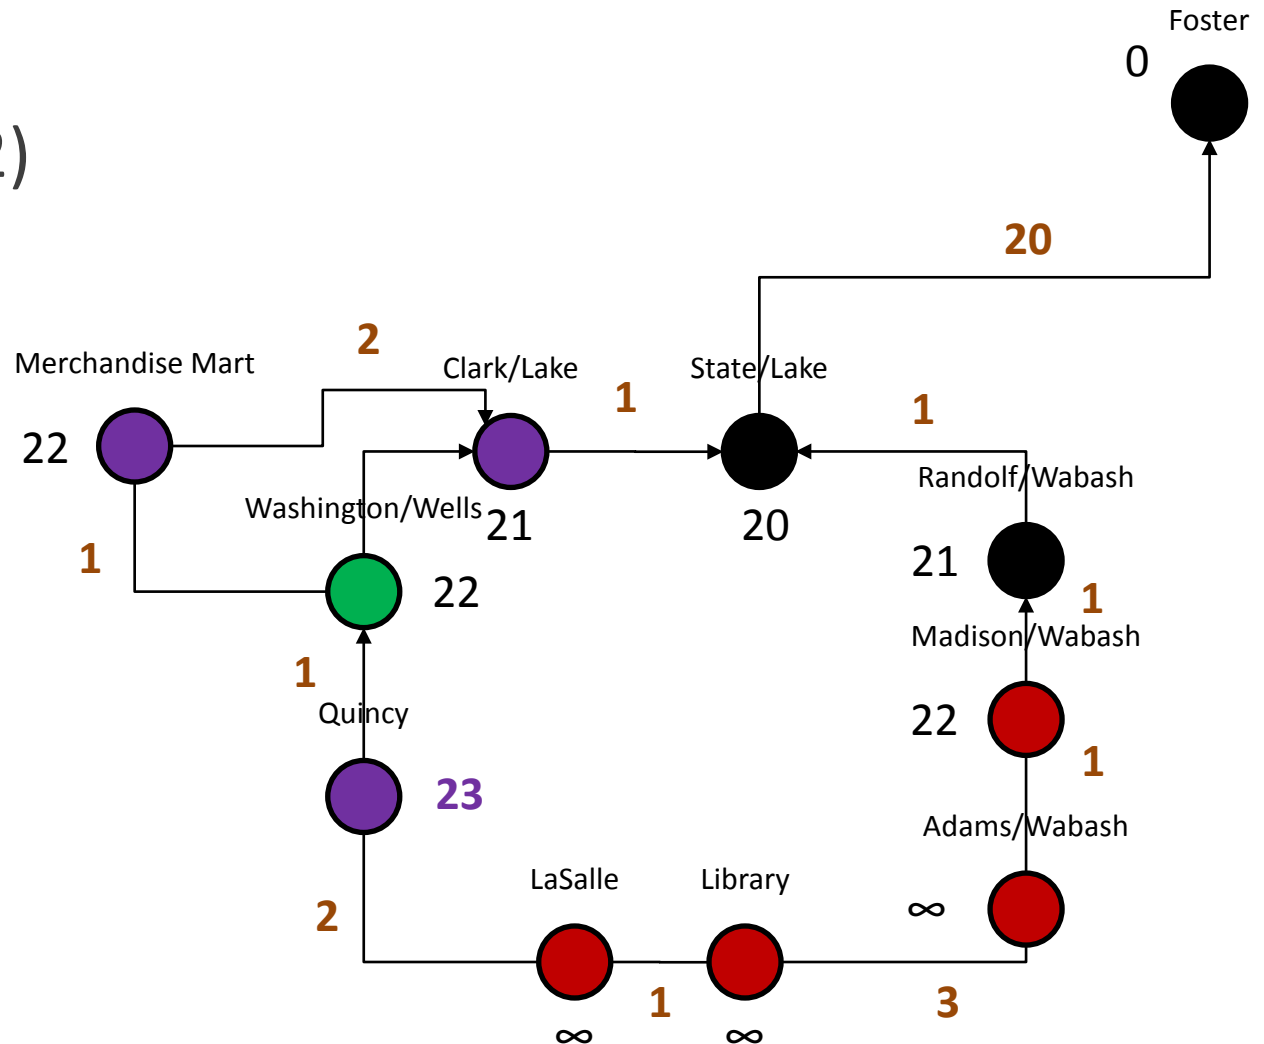
- Madison (22)
- Adams (∞)
- LaSalle (∞)
- Library (∞)
- Quincy (∞)



Update neighbors

Queue:

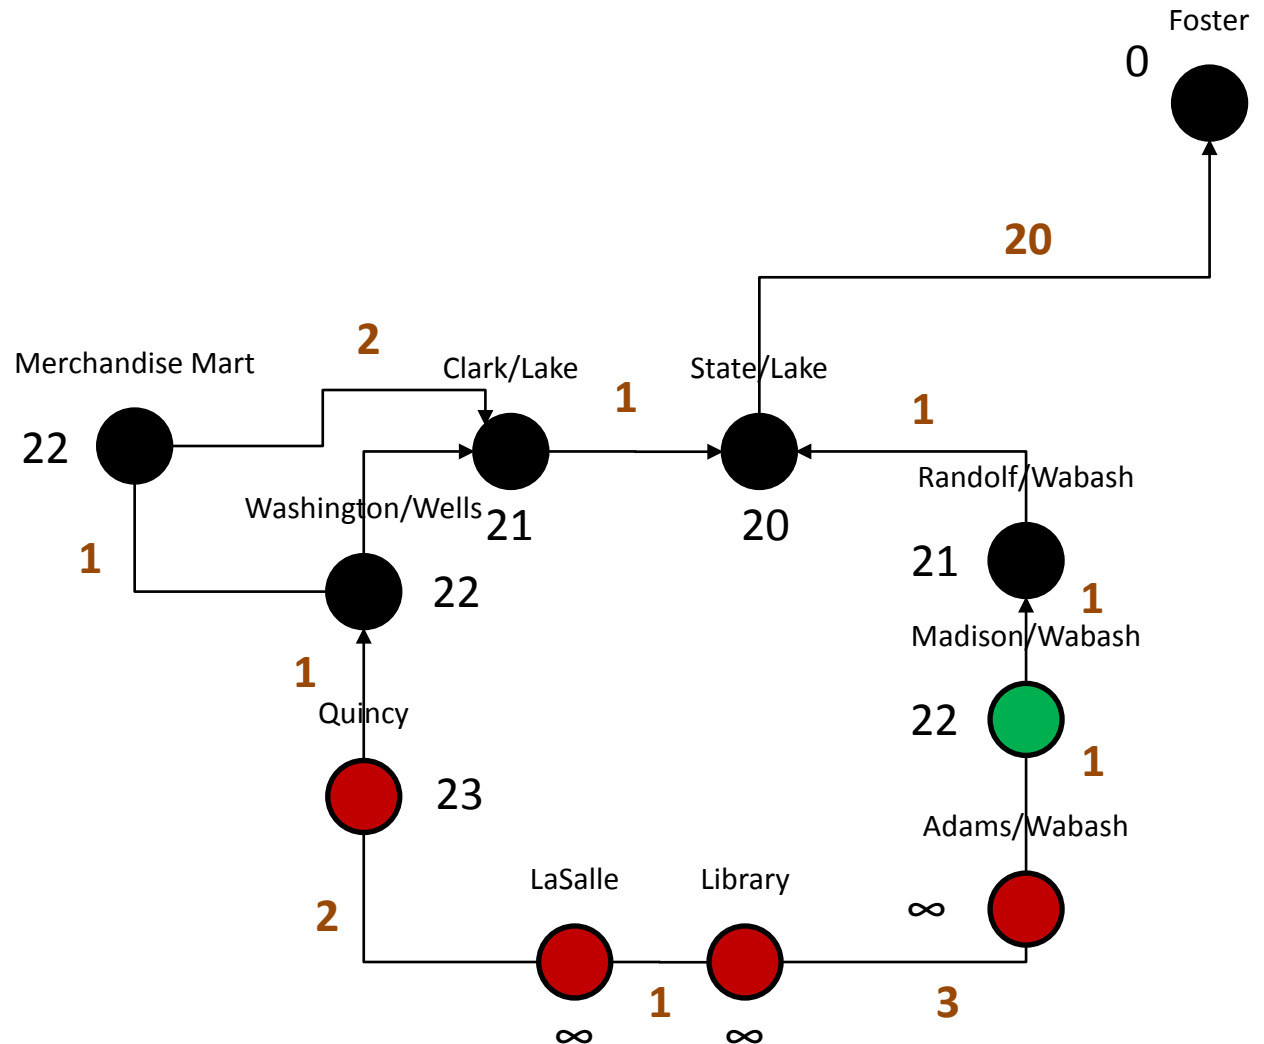
- Madison (22)
- **Quincy (23)**
- Adams (∞)
- LaSalle (∞)
- Library (∞)



Extract min (Madison)

Queue:

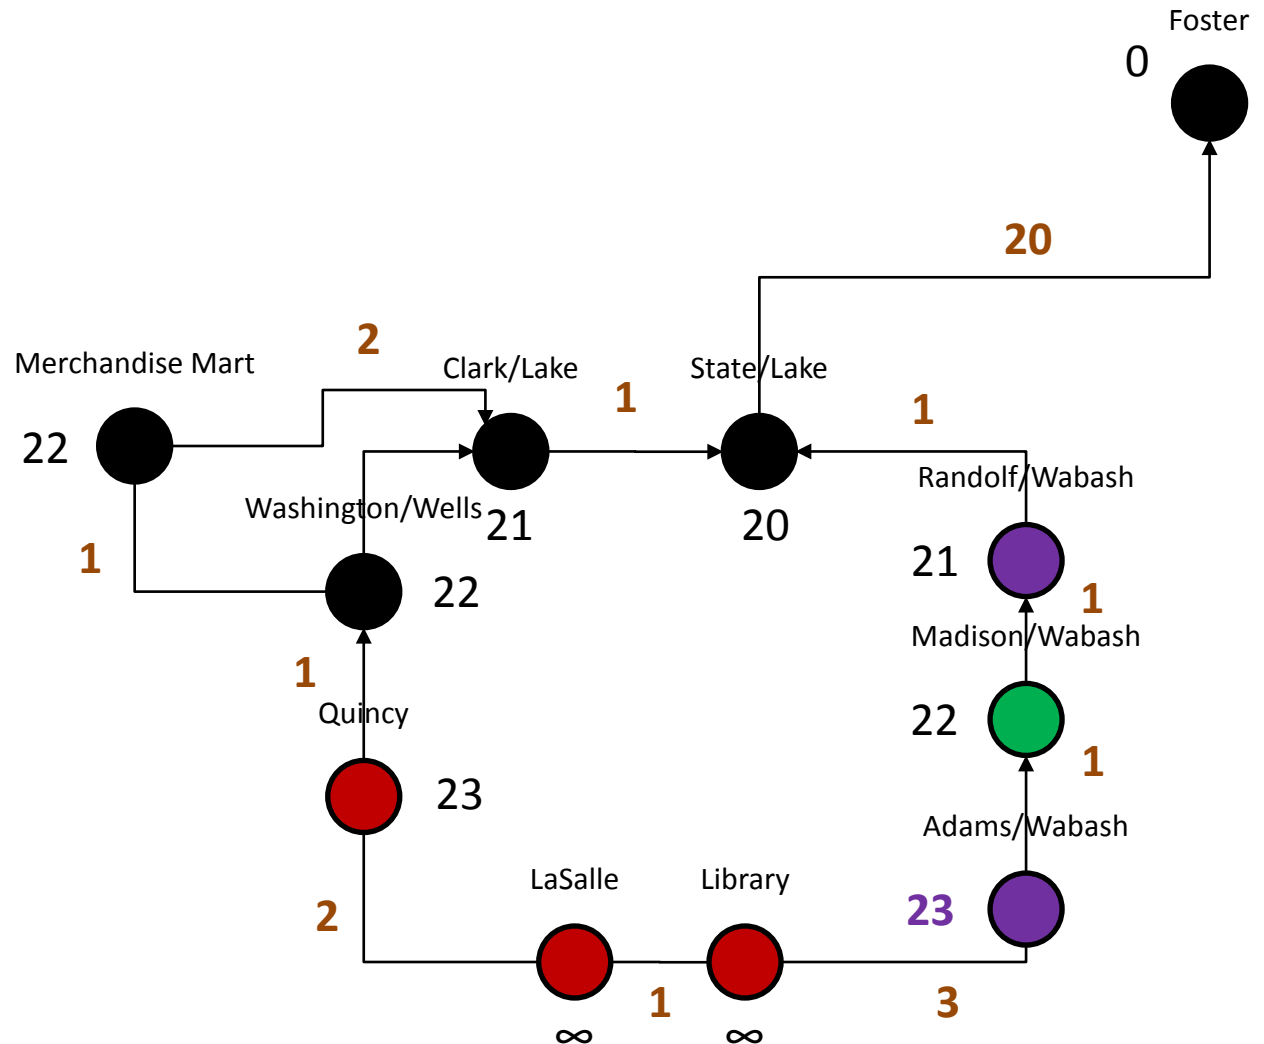
- Quincy (23)
- Adams (∞)
- LaSalle (∞)
- Library (∞)



Update neighbors

Queue:

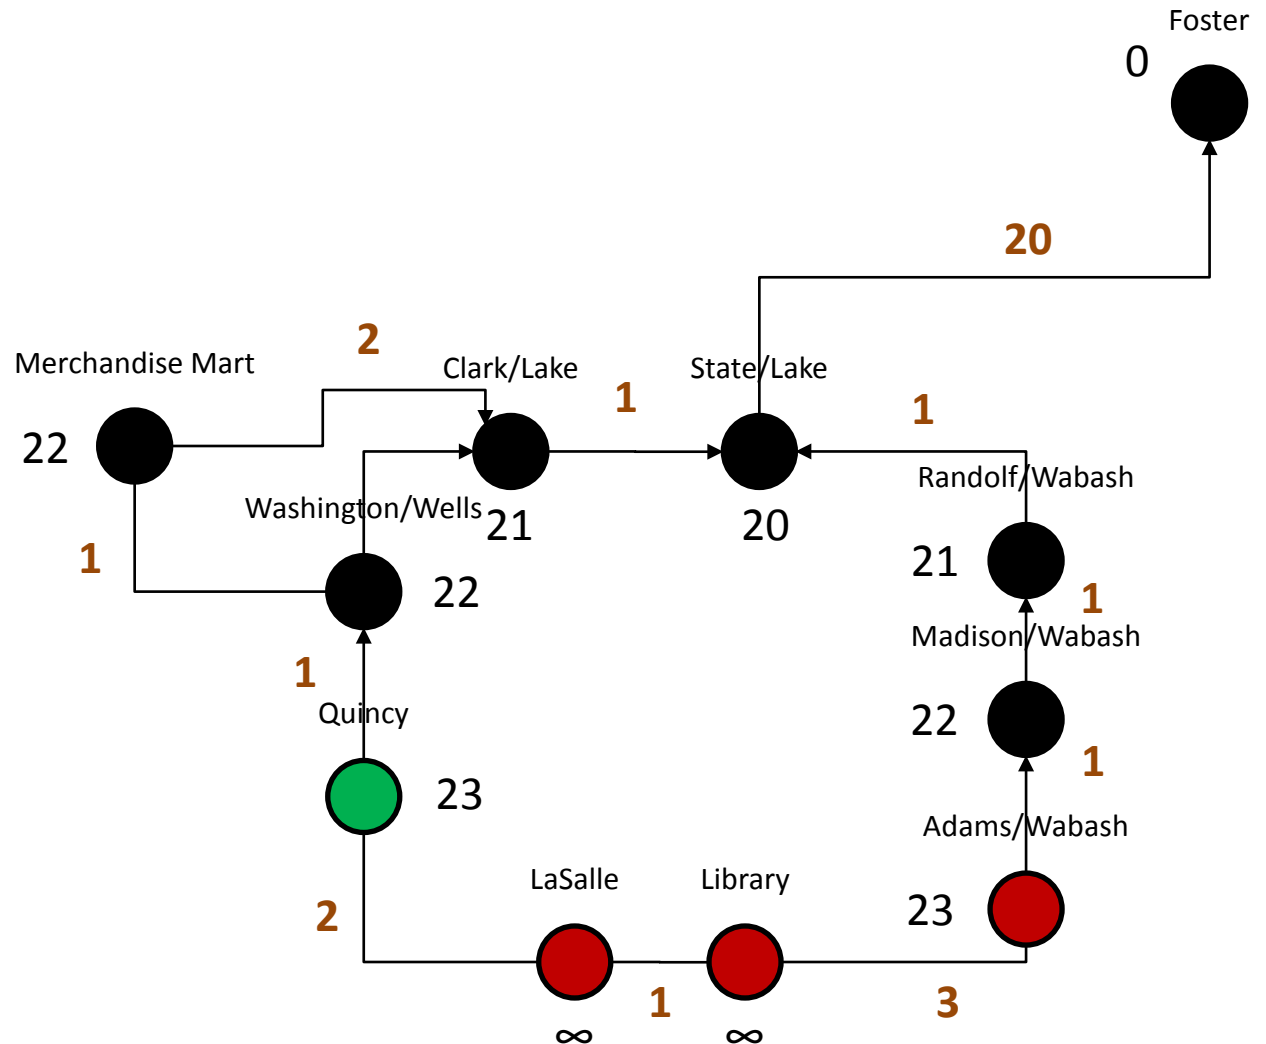
- Quincy (23)
- **Adams (23)**
- LaSalle (∞)
- Library (∞)



Extract min (Quincy)

Queue:


- Adams (23)
- LaSalle (∞)
- Library (∞)



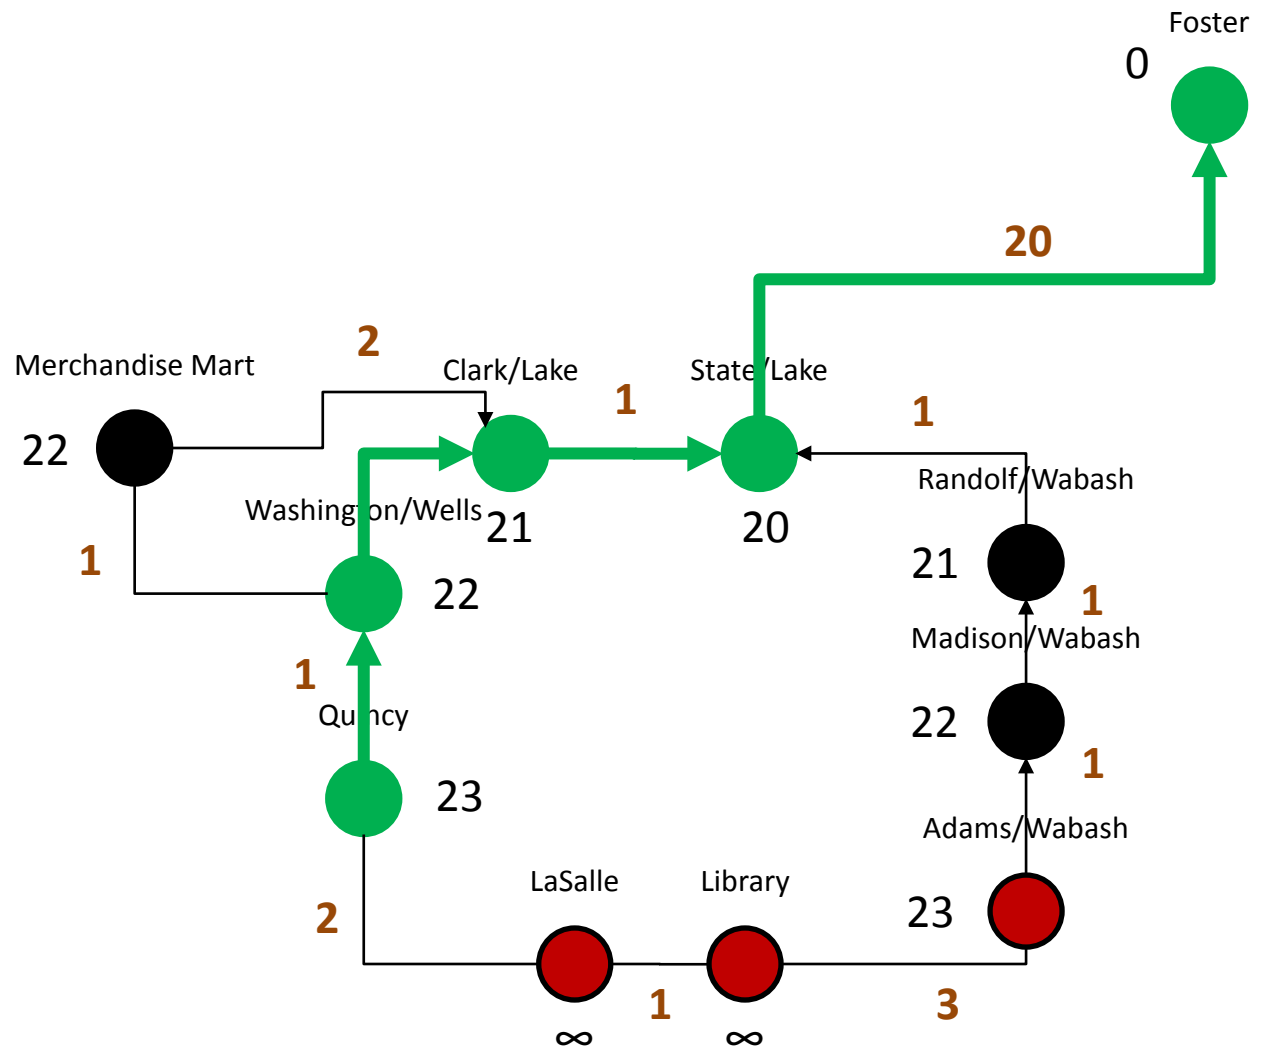
Done!

Queue:

- Adams (23)
- LaSalle (∞)
- Library (∞)



And we have our minimum cost path



How long does it take?

Dijkstra(G, s, e)

PQ = new priority queue

Set all node costs to infinity

s.cost = 0

for each node n in G

 PQ.Insert(n, n.cost)

while PQ not empty

 u = PQ.ExtractMin()

 if u == end then done!

 for each neighbor v of u

 w = weight of edge from u to v

 newCost = u.cost + w

 if newCost < v.cost

 PQ.DecreaseKey(v, newCost)

 v.cost = newCost

 v.predecessor = u

Runs

| Once

| Once

| Once

| $O(V)$ times

| $O(V)$ times

| $O(E)$ times

How long does it take?

Dijkstra(G, s, e)

PQ = new priority queue

Set all node costs to infinity

s.cost = 0

for each node n in G

 PQ.Insert(n, n.cost)

while PQ not empty

 u = PQ.ExtractMin()

 if u == end then done!

 for each neighbor v of u

 w = weight of edge from u to v

 newCost = u.cost + w

 if newCost < v.cost

 PQ.DecreaseKey(v, newCost)

 v.cost = newCost

 v.predecessor = u

Time per execution

| $O(1)$

| $O(V)$

| $O(1)$

| $O(\log V)$

| $O(\log V)$

| $O(1)$

| $O(1)$

| $O(\log V)$

| $O(1)$

How long does it take?

Dijkstra(G, s, e)	Runs	Time	Total
PQ = new priority queue	Once	$O(1)$	$O(1)$
Set all node costs to infinity	Once	$O(V)$	$O(V)$
s.cost = 0	Once	$O(1)$	$O(1)$
for each node n in G	$O(V)$ times	$O(\log V)$	$O(V \log V)$
PQ.Insert(n, n.cost)			
while PQ not empty	$O(V)$ times	$O(\log V)$	$O(V \log V)$
u = PQ.ExtractMin()		$O(1)$	$O(V)$
if u == end then done!			
for each neighbor v of u	$O(E)$ times	$O(1)$	$O(E)$
w = weight of edge from u to v			
newCost = u.cost + w			
if newCost < v.cost		$O(\log V)$	$O(E \log V)$
PQ.DecreaseKey(v, newCost)		$O(1)$	$O(E)$
v.cost = newCost			
v.predecessor = u			

How long does it take?

$$\begin{aligned} &O(1) + O(V) + O(E) \\ &\quad + O(V \log V) + O(E \log V) \\ &= O(1) + O(V + E) + O((V + E) \log V) \\ &= O(V + E) + O((V + E) \log V) \\ &= O((V + E)(1 + \log V)) \\ &= O((V + E) \log V) \end{aligned}$$