

Lecture 17


Dynamic set partitions and the union-find algorithm

EECS-214

Set partitions

- **Division** of elements of a set into **disjoint groups**
- Formally, given a set S , a partition $P = \{ P_1, \dots, P_n \}$ is
 - A set of non-empty subsets of S
 - i.e. for all i , $\{ \} \neq P_i \subseteq S$
 - Such that
 - Every element of S is a member of **exactly one element** of P
 - Or equivalently, the P_i are
 - **Disjoint**


$$S = \bigcup_{P_i \in P} P_i$$



Set partitions

- **Division** of elements of a set into **disjoint groups**
- Formally, given a set S , a partition $P = \{ P_1, \dots, P_n \}$ is
 - A set of non-empty subsets of S
 - i.e. for all i , $\{ \} \neq P_i \subseteq S$
 - Such that
 - Every element of S is a member of **exactly one element** of P
 - Or equivalently, the P_i are
 - **Disjoint**
 - **Cover** S


$$S = \bigcup_{P_i \in P} P_i$$



Set partitions

- **Division** of elements of a set into **disjoint groups**
- Formally, given a set S , a partition $P = \{ P_1, \dots, P_n \}$ is
 - A set of non-empty subsets of S
 - i.e. for all i , $\{ \} \neq P_i \subseteq S$
 - Such that
 - Every element of S is a member of **exactly one element** of P
 - Or equivalently, the P_i are
 - **Disjoint**
 - **Cover** S


$$S = \bigcup_{P_i \in P} P_i$$



Set partitions

- **Division** of elements of a set into **disjoint groups**
- Formally, given a set S , a partition $P = \{ P_1, \dots, P_n \}$ is
 - A set of non-empty subsets of S
 - i.e. for all i , $\{ \} \neq P_i \subseteq S$
 - Such that
 - Every element of S is a member of **exactly one element** of P
 - Or equivalently, the P_i are
 - **Disjoint**
 - **Cover** S

$$S = \bigcup_{P_i \in P} P_i$$



Partitions and equivalence relations

Remember equivalence relations?


- Relation (call it \equiv)
 - Over some set S
- That has the **properties** of
 - **Reflexivity**
 - $x \equiv x$
 - **Transitivity**
 - If $x \equiv y$ and $y \equiv z$
 - Then $x \equiv z$
 - **Symmetry**
 - If $x \equiv y$, then $y \equiv x$
- Divide elements into **equivalence classes**
 - Two elements, $x, y \in S$, are in the same equivalence class iff $x \equiv y$
 - Equivalence class of x is
$$[x] = \{ y \mid x \equiv y \}$$

Partitions and equivalence relations are **interchangeable** ideas

- **Equivalence classes form a partition** of S
- **Given a partition**, we can **construct an equivalence relation** from it
 - $a \equiv b$ iff a and b are in the same set within the partition


Connected components

- **Connectivity** in an undirected graph is an **equivalence relation**
 - $a \equiv b$ iff there is a **path** between a and b
- Its **equivalence classes** are essentially its **connected components**
- They form a **partition** of the nodes
 - Every node belongs to **exactly one connected component**



Connected components

- How do we **compute** the partition?



Connected component analysis using depth-first search

LabelConnectedComponents()

component = 0

for each node in graph

if node not visited

LCCVisit(node, **component**)

component++

LCCVisit(node, c)


mark node visited

node.component = c

for each unvisited

neighbor, n, of node

LCCVisit(n, **c**)



Connected component analysis using depth-first search

LabelConnectedComponents()

component = 0

for each node in graph

if node not visited

LCCVisit(node, component)

component++

LCCVisit(node, c)


mark node visited

node.component = c

for each unvisited

neighbor, n, of node

LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


mark node visited

node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


mark node visited

node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node

 LCCVisit(n, c)



Connected component analysis using depth-first search

LabelConnectedComponents()

 component = 0

 for each node in graph

 if node not visited

 LCCVisit(node, component)

 component++

LCCVisit(node, c)


 mark node visited

 node.component = c

 for each unvisited

 neighbor, n, of node


 LCCVisit(n, c)



Testing if two nodes are connected


This lets us **test connectivity** in **constant time**

- Just look up the **component numbers** of the two nodes
- And check if they're the **same**




Testing if two nodes are connected

- Oh, **your boss called**
- They want to **add an edge**
- You **don't mind do you?**



Dynamic set partition algorithms


- How do we modify this to allow us to **incrementally merge** equivalence classes?



Dynamic set partition algorithms

Basic idea:

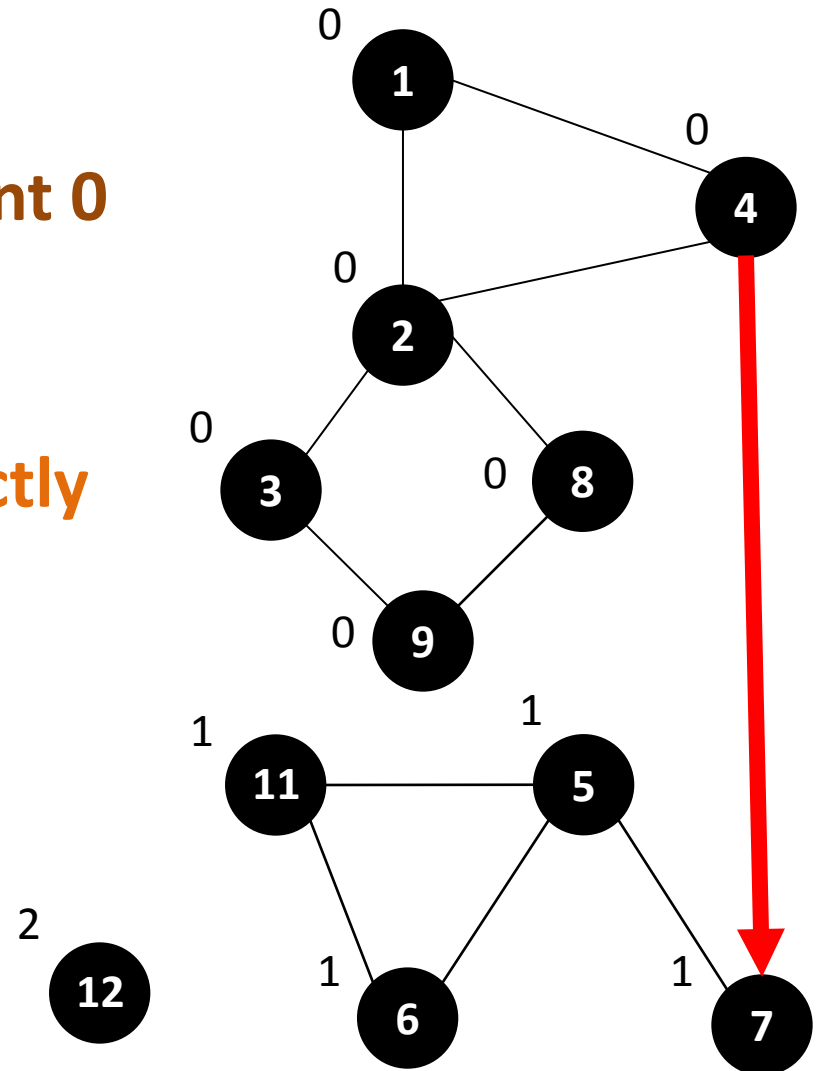
- Remember that **component 0** is “really” component 1



Dynamic set partition algorithms

Basic idea:


- Remember that **component 0** is “really” component 1
- Then **don’t compare** the component numbers **directly**
- Compare the “**real**” component numbers



Dynamic set partition algorithms

Example:

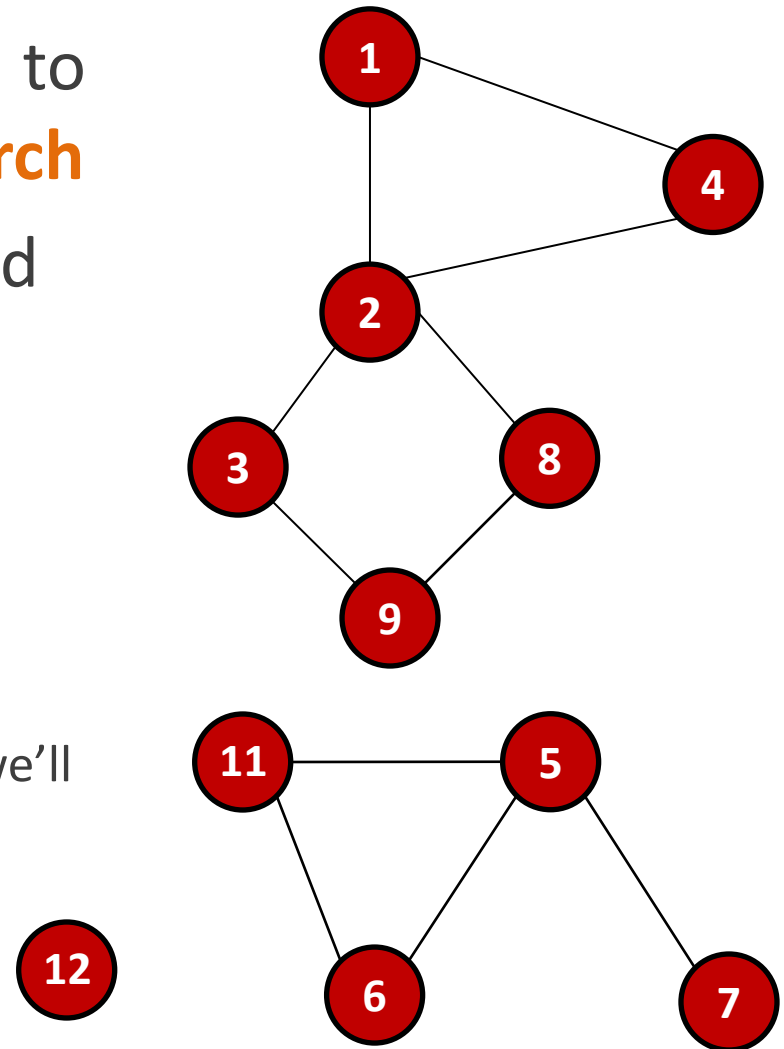
- Node 0 and node 11 are **connected**
 - Because node 0 is in **component 0**
 - But component 0 is **really component 1**
 - And **node 11 is also in component 1**



Computing connected components


- In fact, we can use this idea to **replace the depth-first search**
- And **compute** the connected components **directly**

Note: this is going to be hand-wavy, but we'll make the algorithms precise later



Computing connected components


Here's the idea



Computing connected components

Here's the idea


- Start by assuming **every node** is in its **own** connected **component**



Computing connected components

Here's the idea


- Start by assuming every node is in its own connected component
- **Pick an edge**



Computing connected components


Here's the idea

- Start by assuming every node is in its own connected component
- Pick an edge
- **Merge** the components of the two nodes
 - By remembering that one
 - Is in the component of the other




Computing connected components

- Pick **another edge**




Computing connected components

- Pick another edge
- **Merge** again




Computing connected components

- Now we have that
 - 1 is really in 4's component
 - 3 is really in 2's component




Computing connected components

- **Pick another** edge




Computing connected components

- Pick another edge
- **Merge** again
 - 4 is now in 2's component
 - But 1 is in 4's component
 - So really, 1, 3, and 4 are all in 2's component



Computing connected components


Keep doing this



Computing connected components

Keep doing this


- And we end up with **arrows**
- From all the **nodes in a connected component**



Computing connected components

Keep doing this


- And we end up with **arrows**
- From all the **nodes in a connected component**
- **Flowing to a single node** in each component



Computing connected components

Keep doing this


- And we end up with **arrows**
- From all the **nodes in a connected component**
- **Flowing to a single node** in each component
- These are called the **representatives** of the components



Computing connected components


Each component forms a **tree**

- With arrows pointing **from children to parents**
- With the **representative** as its **root**
- So they collectively form a **forest**



Computing connected components

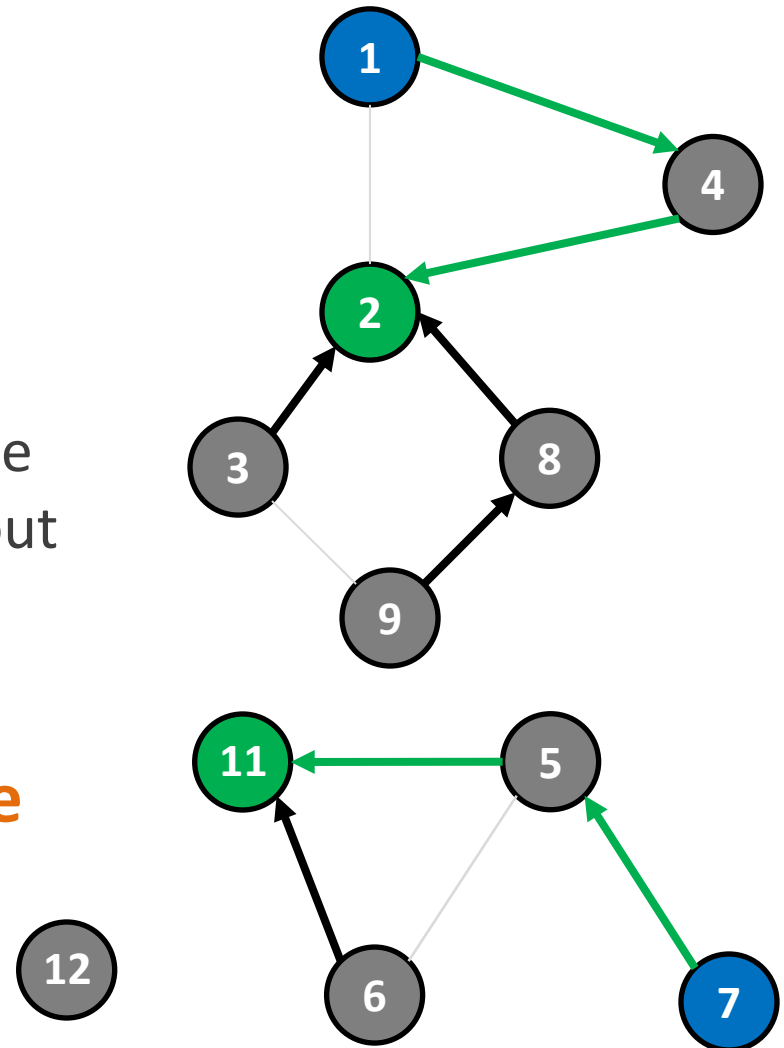
To decide if two nodes are connected



Computing connected components

To decide if two nodes are connected


- **Follow the arrows** to their representatives
 - Easy to do because each node only has one edge pointing out
 - Except representatives, who have none
- Check if they have the **same representative**



Computing connected components

The cool thing is


- To **merge two components**



Computing connected components

The cool thing is


- To **merge two components**
- All we have to do is **add a pointer**
 - From one component's representative
 - To the other's



Computing connected components


The cool thing is

- To **merge two components**
- All we have to do is **add a pointer**
 - From one component's representative
 - To the other's
- Now they both have the **same representative**



The union-find algorithm


- This is a handwavy version of a **general algorithm**
- Called the **union-find** algorithm
- For working with **partitions** of sets



The union-find algorithm


```
class SetElement {  
    SetElement parent;  
}
```

- Keep a **parent pointer** for each element of the set
 - Points to its parent in the tree
 - **Representatives** are roots, so parent is **null**
 - Or it's often **set to itself**
 - Which saves some work later




The union-find algorithm

- We start with all the elements in their **own groups** in the partition




The union-find algorithm

- We start with all the elements in their **own groups** in the partition
- So they're all their **own representatives**



The union-find algorithm


To **union** (merge) two groups together



The union-find algorithm

To **union** (merge) two groups together


- Just make the **representative of one**
- Be the **parent** of the **representative** of the other



The union-find algorithm

To **union** (merge) two groups together


- Just make the **representative of one**
- Be the **parent** of the **representative** of the other



The union-find algorithm

To **union** (merge) two groups together


- Just make the **representative of one**
- Be the **parent** of the **representative** of the other



The union-find algorithm

To **union** (merge) two groups together


- Just make the **representative of one**
- Be the **parent** of the **representative** of the other



The union-find algorithm

To **union** (merge) two groups together


- Just make the **representative of one**
- Be the **parent** of the **representative** of the other



The union-find algorithm

To **union** (merge) two groups together

- Just make the **representative of one**
- Be the **parent** of the **representative** of the other



The union-find algorithm

(slow version)

**// Find the representative
// for e's group**

```
Find(SetElement e) {  
    while (e.parent != null)  
        e = e.parent;  
    return e;  
}
```

**// Merge the groups of a
// and b together**

```
Union(SetElement a, SetElement b)  
{  
    ar = Find(a);  
    br = Find(b);  
    if (ar != br)  
        ar.parent = br;  
}
```

Complexity analysis

// Find the representative
// for e's group

```
Find(SetElement e) {  
    while (e.parent != null)  
        e = e.parent;  
    return e;  
}
```

// Merge the groups of a
// and b together

```
Union(SetElement a, SetElement b)  
{  
    ar = Find(a);  
    br = Find(b);  
    if (ar != br)  
        ar.parent = br;  
}
```

Complexity analysis

// Find the representative
// for e's group

```
Find(SetElement e) {  
    while (e.parent != null)  
        e = e.parent;  
    return e;  
}
```

$O(h)$

// Merge the groups of a
// and b together

```
Union(SetElement a, SetElement b)  
{  
    ar = Find(a);  
    br = Find(b);  
    if (ar != br)  
        ar.parent = br;  
}
```

$O(h)$

The worst-case running time of each algorithm is the height of the tree

Complexity analysis

// Find the representative
// for e's group

```
Find(SetElement e) {  
    while (e.parent != null)  
        e = e.parent;  
    return e;  
}
```

$O(n)$

// Merge the groups of a
// and b together

```
Union(SetElement a, SetElement b)  
{  
    ar = Find(a);  
    br = Find(b);  
    if (ar != br)  
        ar.parent = br;  
}
```

$O(n)$

But the worst-case height of the tree is the number of nodes :-)

How do we make it better?

How do we make it better?

- **Balance** the tree!
- Keep **track** of the **heights** of the trees
- Always **add the shorter one to the bigger one**

Union by rank (height)

```
class SetElement {  
    SetElement parent;  
    int rank;  
}
```

```
Union(SetElement a, SetElement b)  
{  
    ar = Find(a);  
    br = Find(b);  
  
    if (ar == br) return;  
  
    if (ar.rank < br.rank)  
        ar.parent = br;  
    else if (ar.rank > br.rank)  
        br.parent = ar;  
    else {  
        br.parent = ar;  
        ar.rank = ar.rank + 1;  
    }  
}
```

Complexity

- Union by rank guarantees the trees are **balanced**
- So their **heights** are $O(\log n)$
- And so the **complexities** are Union and Find are also $O(\log n)$

But it turns out we can actually do **better**

Saving redundant work

- If we call **Find** on the same node **twice**
- We run up the chain twice
- Why not **cache** the original result?

Path compression

- Just **update** the **parent** of the element
- To point directly at the **representative**
- And while you're at it, update **all the nodes on the path** to the representative too!

```
Find(SetElement a) {  
    if (a.parent == null)  
        return a;  
    a.parent = Find(a.parent);  
    return a.parent;  
}
```

Complexity analysis

- The bad news is that Find is **still $O(\log n)$** in the **worst case**
- So what's the good news?

Amortized complexity analysis

- It's only $O(\log n)$ the **first time** you do a lookup on a node
- After that, it's **constant time**!
 - At least **until** you do **another merge**
- Tarjan proved that the **amortized complexity** of both union and find is **$O(\alpha(n))$**
- Great! **What's $\alpha(n, n)$?**

The Ackermann function

There's this function called the **Ackermann function**:

$$A(m, n) = \begin{cases} n + 1, & \text{if } m = 0 \\ A(m - 1, 1), & \text{if } m > 0 \text{ and } n = 0 \\ A(m - 1, A(m, n - 1)), & \text{otherwise} \end{cases}$$

The Ackermann function

Here's Wikipedia's table of **values** for Ackermann's function:

Values of $A(m, n)$						
$m \backslash n$	0	1	2	3	4	n
0	1	2	3	4	5	$n + 1$
1	2	3	4	5	6	$n + 2 = 2 + (n + 3) - 3$
2	3	5	7	9	11	$2n + 3 = 2 \cdot (n + 3) - 3$
3	5	13	29	61	125	$2^{(n+3)} - 3$
4	13 $= 2^{2^2} - 3$	65533 $= 2^{2^{2^2}} - 3$	$2^{65536} - 3$ $= 2^{2^{2^{2^2}}} - 3$	$2^{2^{65536}} - 3$ $= 2^{2^{2^{2^{2^2}}}} - 3$	$2^{2^{2^{65536}}} - 3$ $= 2^{2^{2^{2^{2^{2^2}}}}} - 3$	$\underbrace{2^{2^{\dots^2}}}_{n+3} - 3$

It grows really really **fast**.

The inverse Ackermann function

- $\alpha(n)$ is the **inverse** of $A(n, n)$
- It grows really really **slowly**

Values of $A(m, n)$

$m \backslash n$	0	1	2	3	4	n
0	1	2	3	4	5	$n + 1$
1	2	3	4	5	6	$n + 2 = 2 + (n + 3) - 3$
2	3	5	7	9	11	$2n + 3 = 2 \cdot (n + 3) - 3$
3	5	13	29	61	125	$2^{(n+3)} - 3$
4	13 $= 2^{2^2} - 3$	65533 $= 2^{2^{2^2}} - 3$	$2^{65536} - 3$ $= 2^{2^{2^{2^2}}} - 3$	$2^{2^{65536}} - 3$ $= 2^{2^{2^{2^{2^2}}}} - 3$	$2^{2^{2^{65536}}} - 3$ $= 2^{2^{2^{2^{2^{2^2}}}}} - 3$	$\underbrace{2^{2^{\dots^2}}}_{n+3} - 3$

The inverse Ackermann function

- It grows **comically slowly**
- α (the number of subatomic particles in the universe) < 4

Values of $A(m, n)$

$m \backslash n$	0	1	2	3	4	n
0	1	2	3	4	5	$n + 1$
1	2	3	4	5	6	$n + 2 = 2 + (n + 3) - 3$
2	3	5	7	9	11	$2n + 3 = 2 \cdot (n + 3) - 3$
3	5	13	29	61	125	$2^{(n+3)} - 3$
4	13 $= 2^{2^2} - 3$	65533 $= 2^{2^{2^2}} - 3$	$2^{65536} - 3$ $= 2^{2^{2^{2^2}}} - 3$	$2^{2^{65536}} - 3$ $= 2^{2^{2^{2^{2^2}}}} - 3$	$2^{2^{2^{65536}}} - 3$ $= 2^{2^{2^{2^{2^{2^2}}}}} - 3$	$\underbrace{2^{2^{\dots^2}}}_{n+3} - 3$

Estimated number of subatomic particles $\cong 10^{80} \cong 2^{265}$

The inverse Ackermann function

- Did I mention **it grows slowly**?

Values of $A(m, n)$

$m \backslash n$	0	1	2	3	4	n
0	1	2	3	4	5	$n + 1$
1	2	3	4	5	6	$n + 2 = 2 + (n + 3) - 3$
2	3	5	7	9	11	$2n + 3 = 2 \cdot (n + 3) - 3$
3	5	13	29	61	125	$2^{(n+3)} - 3$
4	13 $= 2^{2^2} - 3$	65533 $= 2^{2^{2^2}} - 3$	$2^{65536} - 3$ $= 2^{2^{2^{2^2}}} - 3$	$2^{2^{65536}} - 3$ $= 2^{2^{2^{2^{2^2}}}} - 3$	$2^{2^{2^{65536}}} - 3$ $= 2^{2^{2^{2^{2^{2^2}}}}} - 3$	$\underbrace{2^{2^{\dots^2}}}_{n+3} - 3$

The inverse Ackermann function

- So slowly that it **might as well be constant**

Values of $A(m, n)$


$m \backslash n$	0	1	2	3	4	n
0	1	2	3	4	5	$n + 1$
1	2	3	4	5	6	$n + 2 = 2 + (n + 3) - 3$
2	3	5	7	9	11	$2n + 3 = 2 \cdot (n + 3) - 3$
3	5	13	29	61	125	$2^{(n+3)} - 3$
4	13 $= 2^{2^2} - 3$	65533 $= 2^{2^{2^2}} - 3$	$2^{65536} - 3$ $= 2^{2^{2^{2^2}}} - 3$	$2^{2^{65536}} - 3$ $= 2^{2^{2^{2^{2^2}}}} - 3$	$2^{2^{2^{65536}}} - 3$ $= 2^{2^{2^{2^{2^{2^2}}}}} - 3$	$\underbrace{2^{2^{\dots^2}}}_{n+3} - 3$

Amortized complexity analysis

- Tarjan proved that the **amortized complexity** of both union and find is $O(\alpha(n))$
- So for all practical purposes, it's **$O(1)$**
- That's pretty **cool**


Incremental connected components

- Add parent and rank fields to each node
- For each edge (u,v)
 $\text{Union}(u, v)$
- To test whether two nodes are connected check
 $\text{Find}(u) == \text{Find}(v)$




Amortized complexity

- For each edge (u,v)
Union(u, v) $O(\alpha(V)E) \cong O(E)$
- To test whether two nodes are connected check
Find(u) == Find(v)
 $O(\alpha(V)) \cong O(1)$



Amortized complexity

- For each edge (u,v)
Union(u, v) $O(\alpha(V)E) \cong O(E)$
- To test whether two nodes are connected check
Find(u) == Find(v)
 $O(\alpha(V)) \cong O(1)$
- So we can compute connected component **incrementally** and still do it in **linear time**



Oh crap! I don't understand the Ackermann function at all! Do I need to know this for the quiz??!?!?

- No
- You just need to understand
 - The **union-find algorithm** with union by **rank** and **path compression**
 - That it's **effectively $O(1)$** amortized time
 - **Why** you might want to keep track of **partitions** on a set
 - E.g. to compute connected components or equivalence classes

Reading

- Read Chapter 21, sections 1-3